

# Tensilica Vision DSP Introduction

*Imaging/Vision BU*  
7<sup>th</sup> Dec 2017

**CONFIDENTIAL**

**cā dence<sup>®</sup>**

# CNN Algorithm Development Trends

Increasing Computational Requirements  
(~16X in <4 years)

- AlexNet (2012)
- Inception (2015)
- ResNet (2015)

| NETWORK      | MACS/IMAGE     |
|--------------|----------------|
| ALEXNET      | 724,406,816    |
| INCEPTION V3 | 5,713,232,480  |
| RESNET-101   | 7,570,194,432  |
| RESNET-152   | 11,282,415,616 |

Network Architectures Changing Regularly

- AlexNet (bigger convolution); Inception V3 and ResNet (smaller convolution)
- Linear network vs. branch

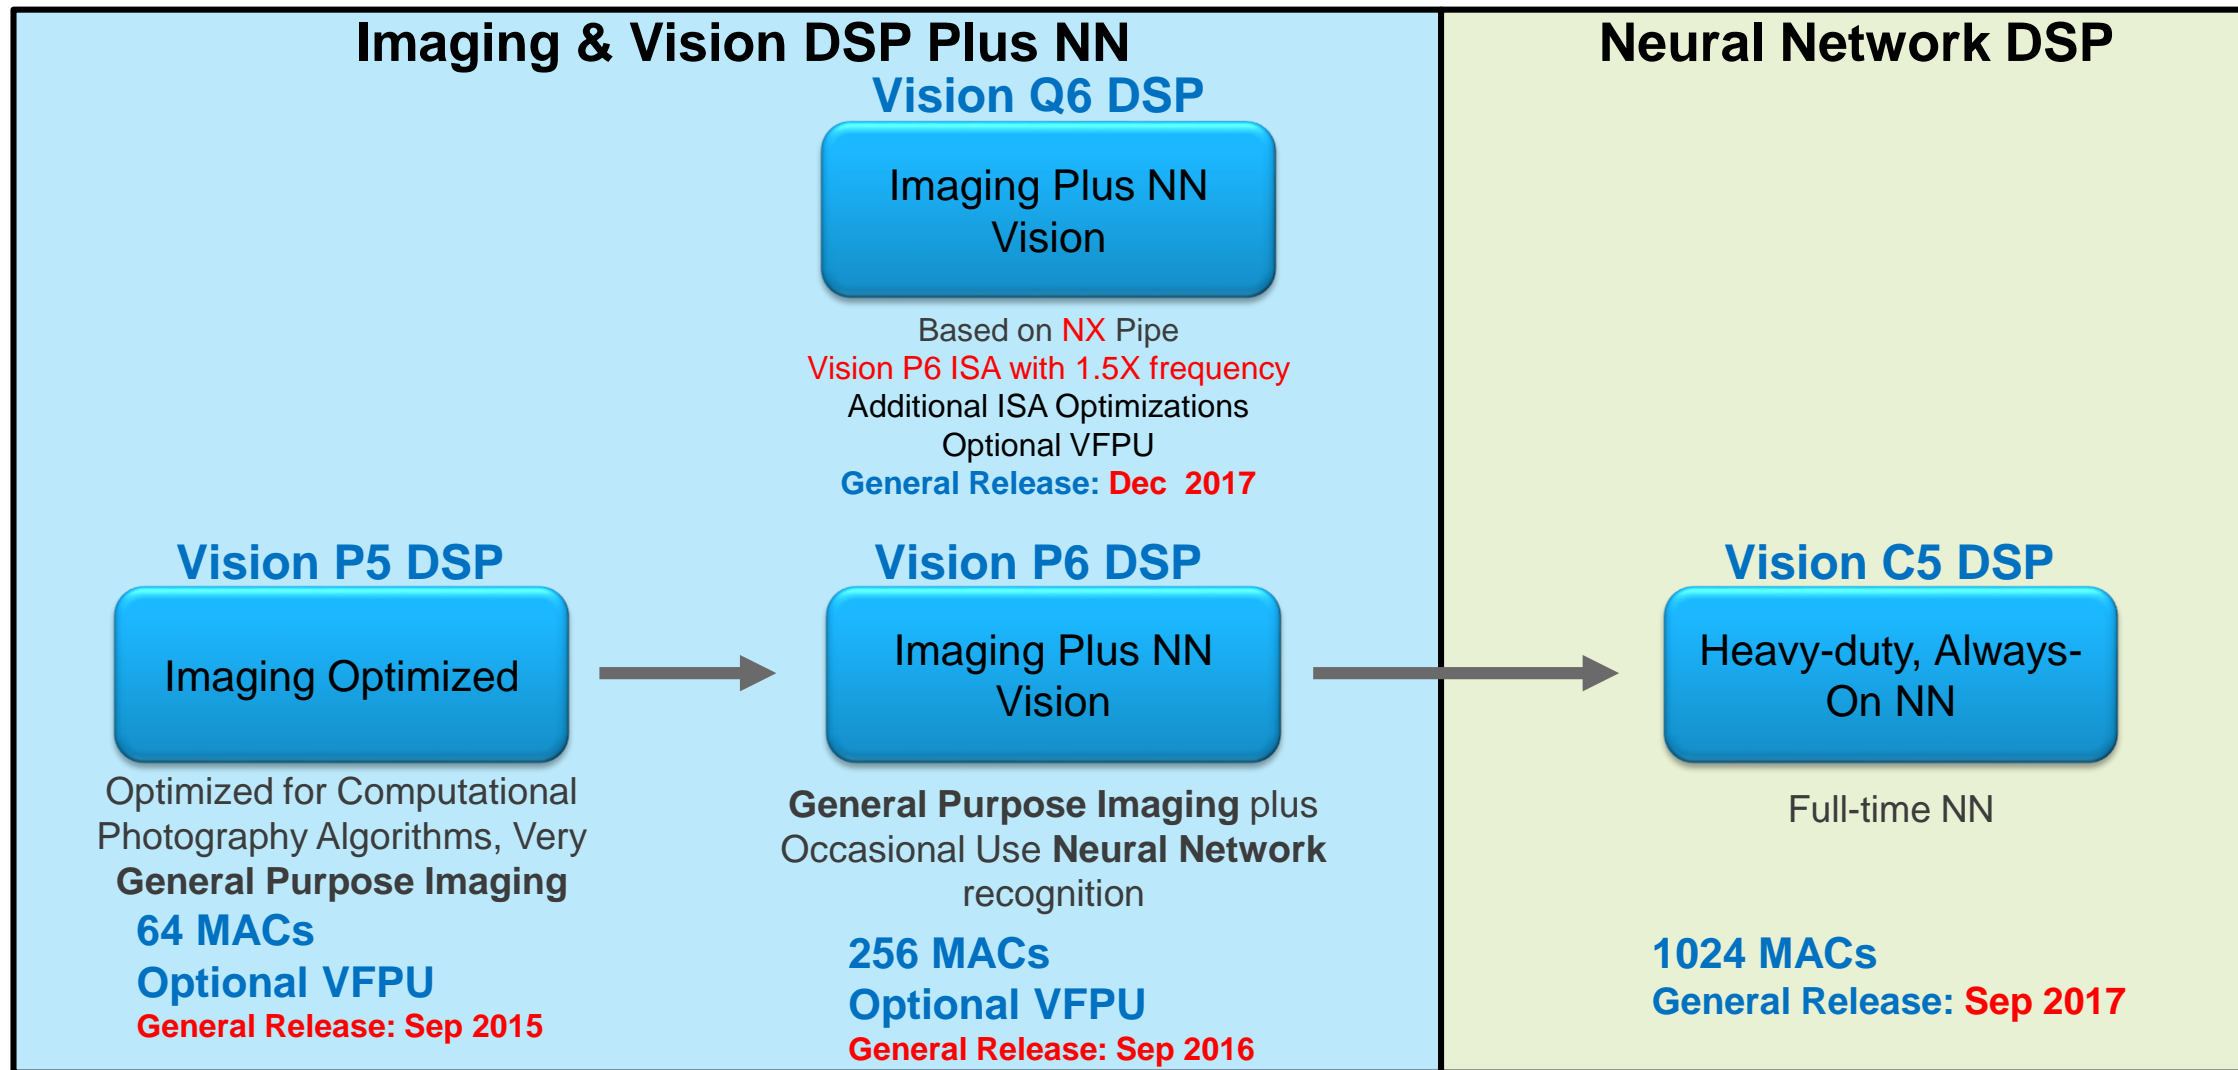
New Applications and Markets

- Automotive, server, home (voice-activated digital assistants), mobile, surveillance

How do you pick an inference hardware platform today (2017) for a product shipping in 2019-2020+? How do you achieve low-power efficiency yet be flexible?

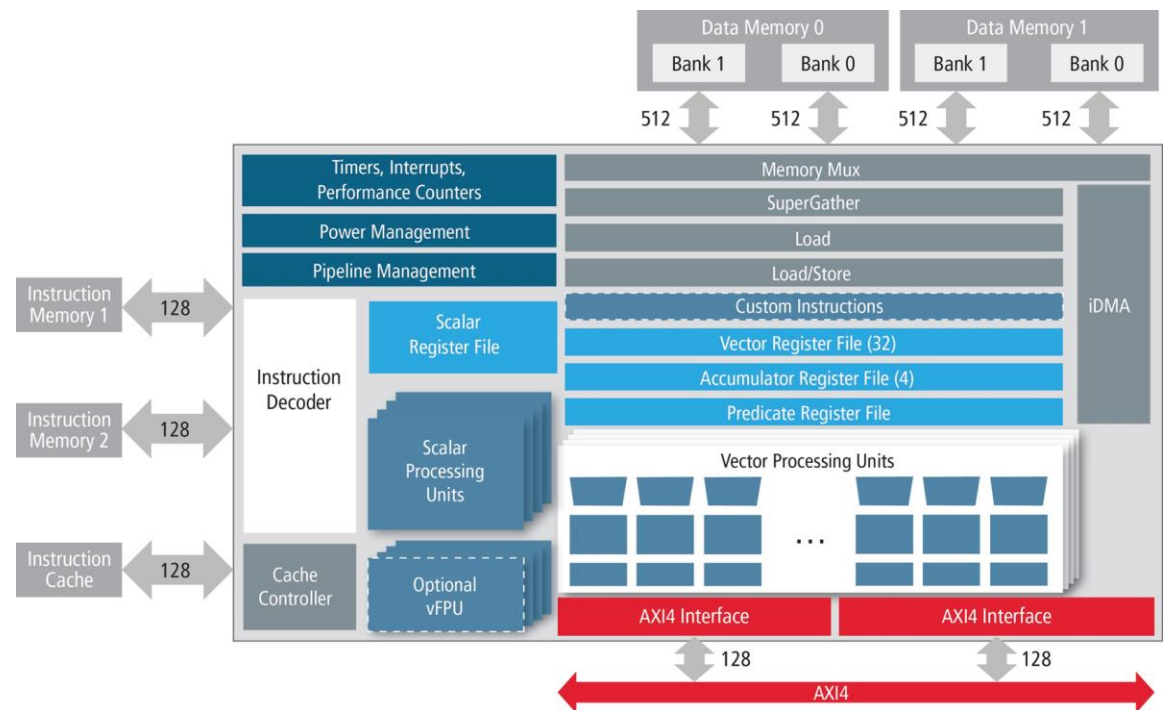
L  
o  
w  
P  
o  
w  
e  
r

# Vision DSPs: Family of Imaging and NN DSPs



Two product lines: Imaging/Vision & Neural network DSP

# Vision P5/P6 Architecture



|                                   |   |
|-----------------------------------|---|
| VLIW & SIMD                       | 5 slots<br>64way 8-bit<br>32way 16-bit<br>16way 32-bit                    |
| ALU Ops<br>(MAX 4 out of 5 slots) | 256 8-bit<br>128 16-bit<br>64 32-bit                                      |
| MAC<br>(1 of 5 slots)             | <b>Vision P5: 64 (8x8)</b><br><b>Vision P6: 256 (8x8)</b>                 |
| Memory Width                      | 1024-bits<br>2 vector load/store units                                    |
| # of Vector Registers             | 32  |
| SuperGather                       | 32 non-contiguous locations read/ written<br>per instruction              |
| Bus Interface                     | AXI4  |
| iDMA                              | no alignment restrictions, local memory to<br>local memory transfers, ... |
| Target Frequency                  | 800Mhz @28nm<br>1.1 GHz @16nm   |
| Optional                          | Vector Floating Point,<br>ECC   |

# Vision P6 CNN Performance (Hand Code Runs)

| Network                               | Author        | Batch Size | FPS @ 1.1GHz<br>(Vision P6) | DDR Latency<br>(in cycles)               |
|---------------------------------------|---------------|------------|-----------------------------|--|
| Inception V3<br>(299x299x3 Input ROI) | Multicoreware | 1          | 33.4                        | With 0 DDR Latency                       |
|                                       |               |            | 32.7                        | With 100 DDR Latency                     |
| Alexnet<br>(227x227x3 Input ROI)      | Cadence       | 8          | 241                         | With 0 DDR Latency                       |
|                                       |               |            | 235                         | With 100 DDR Latency                     |
| Fast YOLO<br>(Vehicle Detection)      | Multicoreware | 1          | 111                         | 480p as input                            |
| GTSR<br>(32x32 Input ROI)             | Cadence       | 1          | 4598                        | 16x8 MAC are used                        |
| Fast YOLO<br>(People Detection)       | BDTi          | 1          | 91                          | 224x224 as input<br>With 100 DDR Latency |

## Note:

- Inception V3, Alexnet & FastYolo are implemented with 8-bit data and 8-bit coefficient
- Reference memory configuration of 64KB 4-way I\$ and 2x128 Data RAM has been used

# Vision P6 CNN Performance (XNNC Runs)

| Network  | Batch Size | FPS @ 1.1GHz<br>(with 100 cycles<br>of DDR Latency) | Quantization Error<br>(Top-1) |
|--|------------|---|-------------------------------|
| VGG-16   | 1          | 11  | 0.3%*                         |
|  | 8          | 12  |                               |
| VGG-19   | 1          | 9   | 0.5%**                        |
| ResNet-50  | 1          | 36  | 0.9%*                         |
| ResNet-101   | 1          | 22  | 0.2%*                         |
|  | 8          | 23  |                               |
| ResNet-152   | 1          | 14  | 0.7%**                        |
| Caffe Variant of<br>Inception V3<br>(GoogleNet-bn) | 1          | 66  | 0.3%*                         |
|  | 8          | 74  |                               |

\*Tested over 50K Images  
\*\*Tested over 15K Images

- Between XNNC and Manual implementation we see ~20% difference, mostly coming from generic library and automatic code generation
- Reference memory configuration of 64KB 4-way I\$ and 2x128 Data RAM has been used
- XNNC FPS are preliminary pending optimizations



# Quantization: Top-1 Accuracy Loss

Floating Point → 8 Bit Fixed Point Weights  
**Less Than 1% Accuracy drop**

**Top-1 Accuracy in % for  
Networks manually Implemented on VP6**

| Network                                     | Floating Point | 16 bit Fixed Point | 8 bit Fixed Point |
|---|----------------|--------------------|-------------------|
| AlexNet*<br>(Cadence)                       | 62.7%          |                    | 61.6%             |
| Fast YOLO on KITTI: mAP*<br>(Multicoreware) | 63.8%          | 63.8%              | 63.5%             |
| Fast YOLO on MCW: mAP*<br>(Multicoreware)   | 65.9%          | 65.8%              | 65.5%             |
| Inception V3*<br>(Multicoreware)            | 74.0%          |                    | 73.3%             |

\*Tested over 50K ImageNet Dataset

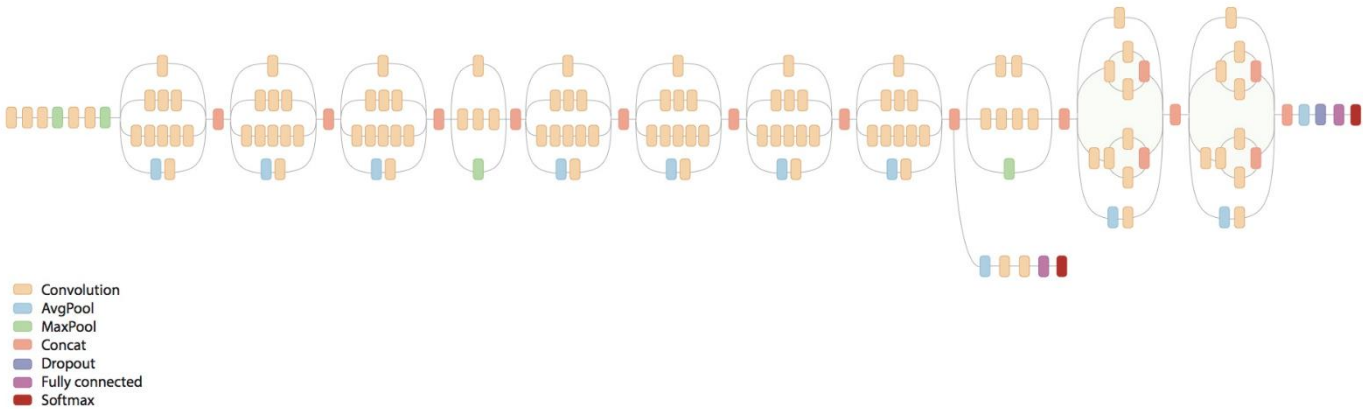
\*\*Tested over 15K ImageNet Dataset

**Top-1 Accuracy in % for  
Networks Implemented on VP6 using XNNC**

| Network                          | Floating Point | 8 bit Fixed Point |
|----------------------------------|----------------|-------------------|
| VGG-16*                          | 68.3%          | 68%               |
| ResNet-101*                      | 71.7%          | 71.5%             |
| Inception V3*<br>(caffe variant) | 70.5%          | 70.2%             |
| ResNet-152**                     | 75.2%          | 74.5%             |
| VGG-19**                         | 67.8%          | 67.4%             |

# Inception V3 Performance on Vision P6 (Batch size of 1) (8bit Data and 8 bit Weights)

| Inception V3 Details                     |           |
|--|-----------|
| Input ROI                                | 299x299x3 |
| Number of Layers                         | 110       |
| Compute Requirement                      | 5.78 GMAC |
| Bandwidth Requirement<br>(8 bit Weights) | 19.4 MB   |



| Vision P6 FPS (@ 1.1GHz)<br>With Tiling and DMA using Reference Configuration<br>(256KB on chip memory) |                         |             |
|---|-------------------------|-------------|
| FPS   | Overall MAC Utilization | DDR Latency |
| 33.4  | 68.61%                  | 0 cycles    |
| 32.7  | 67.16%                  | 100 cycles  |

\*Higher Batch sizes can provide better FPS

| Vision P6 Accuracy (Loss <1%)<br>Using 8bit Quantized Data & Weights |        |                  |
|--|--------|------------------|
| Accuracy*  | Float  | 8bit Fixed Point |
| Top-1 Accuracy   | 74.00% | 73.29%           |
| Top-5 Accuracy   | 91.62% | 91.18%           |

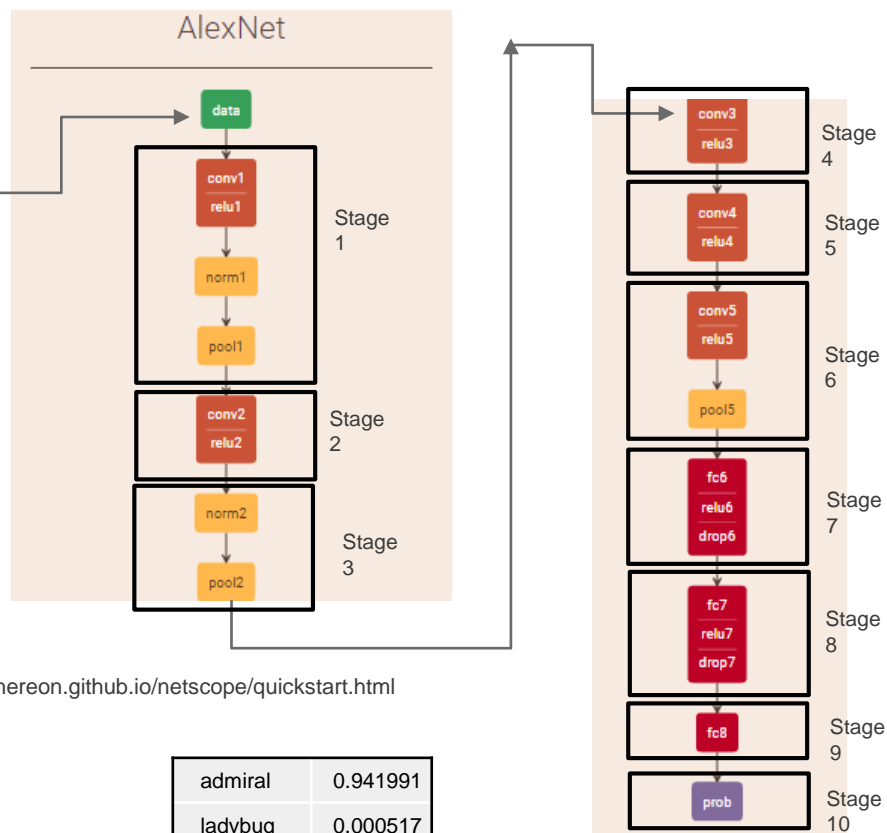
\*Accuracy tested over 50K images in ImageNet Val set



# Vision P6 Running Alexnet Convolutional Neural Network



227x227 RGB image



Alexnet visualization from <http://ethereon.github.io/netscope/quickstart.html>

Vision P6 Detection Result

|          |          |
|----------|----------|
| admiral  | 0.941991 |
| ladybug  | 0.000517 |
| monarch  | 0.000287 |
| tench    | 0.000057 |
| goldfish | 0.000057 |

## Alexnet:

Winner of the ImageNet (ILSVRC) 2012 Contest  
Trained for 1000 different classes (images)  
Most often quoted benchmark for CNN  
Classifier CNN Example  
5 Conv & 3 FC layers  
Input image: 227x227 image patch (ROI)

## Cadence Alexnet Implementation

Based on Caffe 32b floating point Alexnet model  
Use 8 bit coefficients, 8 bit data computations  
Pure C P6 implementation,  
No library dependencies such as BLAS, NumPy, etc

## Performance on Vision P6 @ 1.1GHz

Including tiling and DMA with DDR latency of 0 is 241 fps  
Including tiling and DMA with DDR latency of 100 is 235 fps

## Weighted dynamic power for core in mW/MHz

**0.298 mW/MHz**

TSMC 16FF+ LL, RVT, 0.8V, 500MHz netlist

# Vision P6 Area: TSMC 16FF+ LL 9-Track

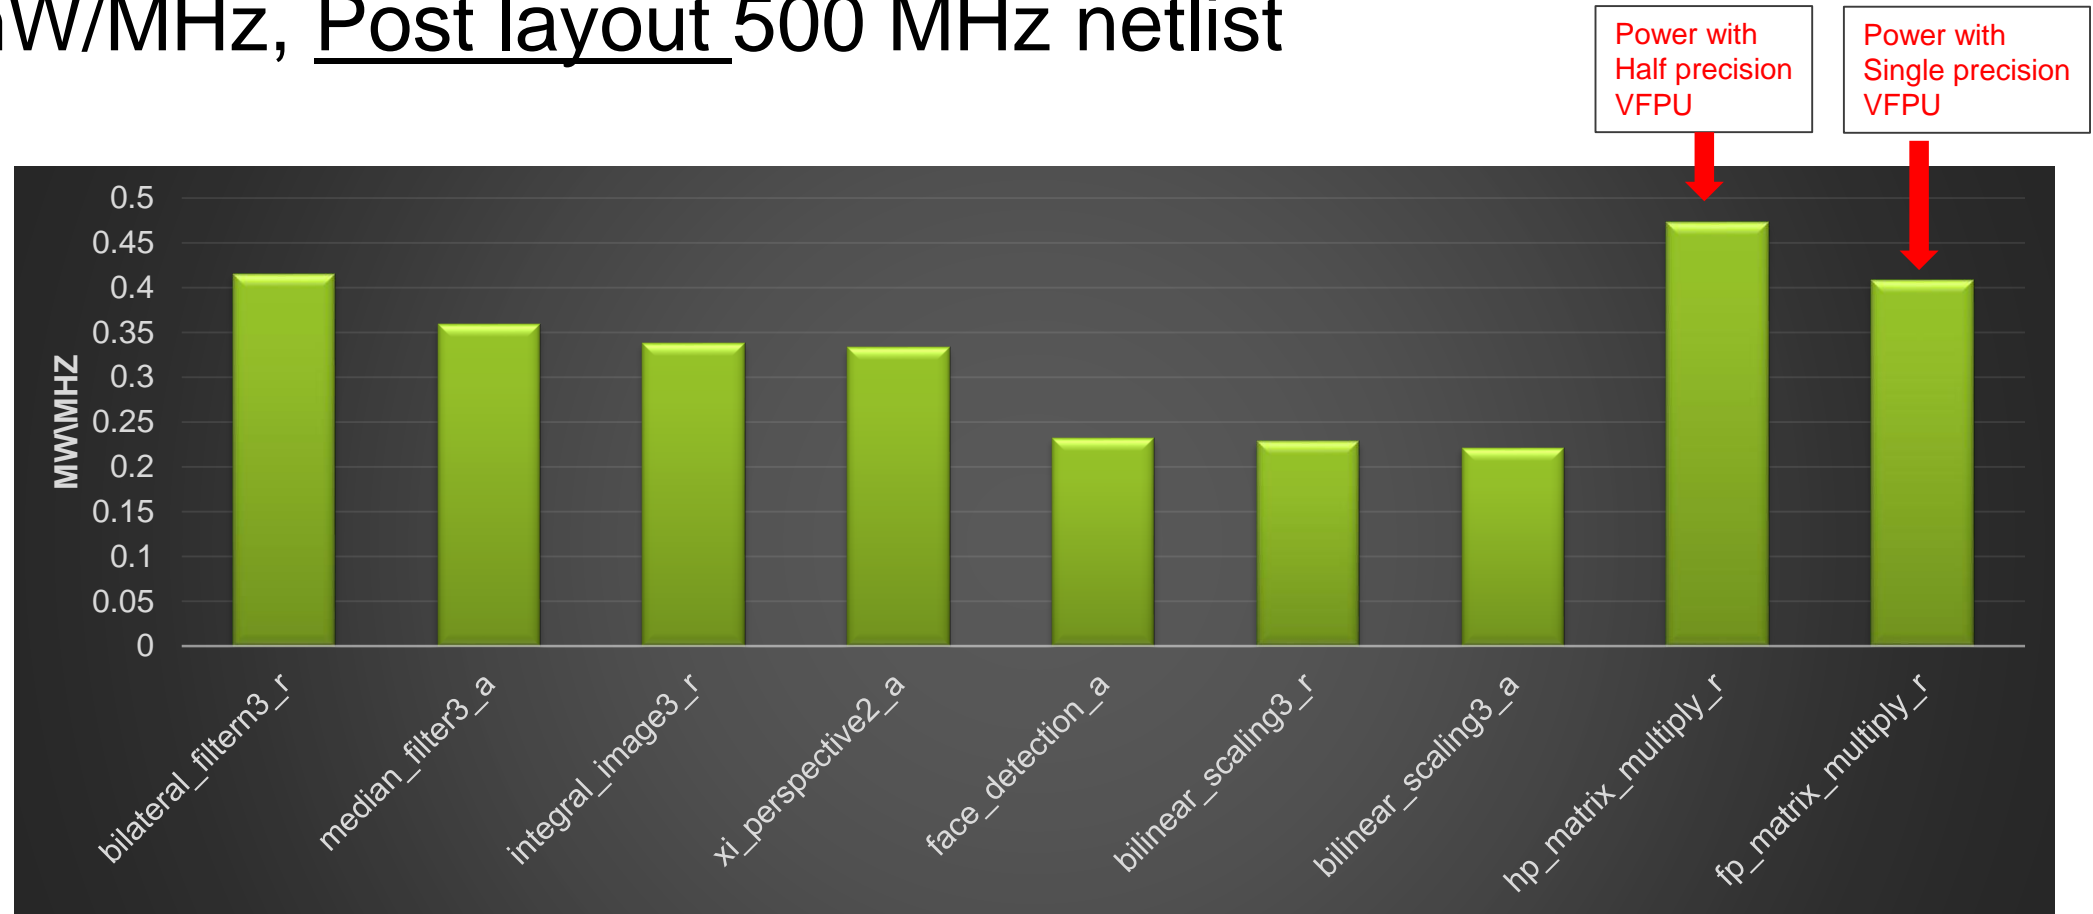
## Post layout Core Cell, Memory Cell and final Floor Plan

| Metric                   | 1.1GHz   | 800 MHz  | 500MHz (VP6P16)  | 500MHz (VP6HP16)  | 500MHz   |   |
|--------------------------|--|--|--|---|--|---|
| Vision P6 reference with | 64KB 4-way I\$<br>2 x 128 KB Data RAMs<br><br><b>No VFPU</b> | 64KB 4-way I\$<br>2 x 128 KB Data RAMs<br><br><b>No VFPU</b> | 64KB 4-way I\$<br>2 x 128 KB Data RAMs<br><br><b>No VFPU</b> | 64KB 4-way I\$<br>2 x 128 KB Data RAMs<br><br><b>Half Precision VFPU (FP16)</b> | 64KB 4-way I\$<br>2 x 128 KB Data RAMs<br><br><b>Single Precision VFPU</b> |   |
|                          | ILVT, LVT and RVT Overdrive library (Nominal VDD=1.0V)       | RVT, LVT, ILVT (VDD=0.8V)                                    | RVT only – Vdd 0.8 V   | RVT only – Vdd 0.8 V  | RVT only – Vdd 0.8 V   |   |
| Memory compiler:         | tcbn16ffplusllbwp16p 90cpd, version 111a                     | ts1n16ffpllvt, version 110b                                  | tcbn16ffplusllbwp16p 90cpd, version 111a                     | tcbn16ffplusllbwp16p 90cpd, version 111a  | tcbn16ffplusllbwp16p 90cpd, version 111a                                   |   |
| Standard Cell Area       | <b>0.58 mm<sup>2</sup></b>                                   | <b>0.576 mm<sup>2</sup></b>                                  | <b>0.531 mm<sup>2</sup></b>                                  | <b>0.615 mm<sup>2</sup></b>   | <b>0.604 mm<sup>2</sup></b>  | Post layout standard cell area of Vision P6 (including iDMA, Supergather) |
| Memory Area              | 0.366 mm <sup>2</sup>  | 0.366 mm <sup>2</sup>  | 0.366 mm <sup>2</sup>  | 0.366 mm <sup>2</sup>   | 0.366 mm <sup>2</sup>  | Area depends on amount of memory, memory compiler                         |

Half Precision VFPU unit adds ~ 15.81% area  
Single Precision VFPU unit adds ~ 13.75% area

# Vision P6 Power: TSMC 16ff+ LL 9 Track

## In mW/MHz, Post layout 500 MHz netlist



Data represent mw/MHz

# Vision P6 Area: TSMC 28HPC+ 9 Track

## Post layout Core Cell, Memory Cell and final Floor Plan

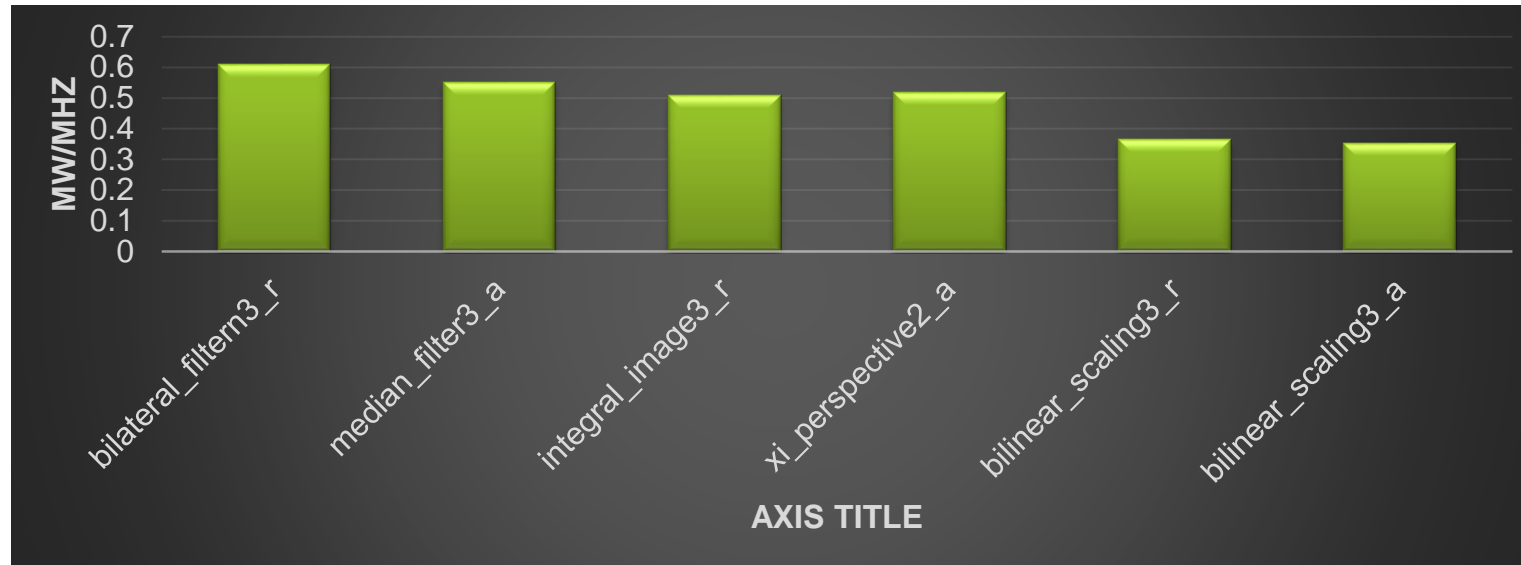
| Metric                     | 500MHz   | 500MHz   | 500MHz (P6V_P28)   |   |
|----------------------------|--|--|--|---|
| Vision P6 reference with   | 64KB 4-way I\$<br>2 x 128 KB Data RAMs<br><br><b>No VFPU</b> | 64KB 4-way I\$<br>2 x 128 KB Data RAMs<br><br><b>Half Precision VFPU</b> | 64KB 4-way I\$<br>2 x 128 KB Data RAMs<br><br><b>Single Precision VFPU</b> |   |
|                            | RVT only –<br>Vdd 0.8 V                                      | RVT only –<br>Vdd 0.8 V  | RVT only –<br>Vdd 0.8 V  |   |
| Memory compiler:           | tcbn28hpcplusbwp35p140,<br>version 110c                      | tcbn28hpcplusbwp35p140,<br>version 110c                                  | tcbn28hpcplusbwp35p140,<br>version 110c                                    |   |
| Post Layout Core Cell Area | <b>1.24 mm<sup>2</sup>*</b>                                  | <b>1.436 mm<sup>2</sup>*</b>   | <b>1.41 mm<sup>2</sup></b>   | Post layout standard cell area<br>of Vision P6 (including iDMA,<br>Supergather) |
| Memory Area                | 0.6 mm <sup>2</sup>  | 0.6 mm <sup>2</sup>  | 0.6 mm <sup>2</sup>  | Area depends on amount of<br>memory, memory compiler                            |

\*Numbers have been derived from VP6+SP  
VFPU based on mentioned scales

Half Precision VFPU unit adds ~ 15.81% area  
Single Precision VFPU unit adds ~ 13.75% area

# Vision P6 Power: TSMC 28HPC+ 9 Track

## In mW/MHz, Post layout 500 MHz netlist



Data represent mw/MHz  
Calculated on VP6 + SP VFPU



# Vision C5: Neural Network DSP



# Tensilica® Vision C5 and Vision P6 DSPs:

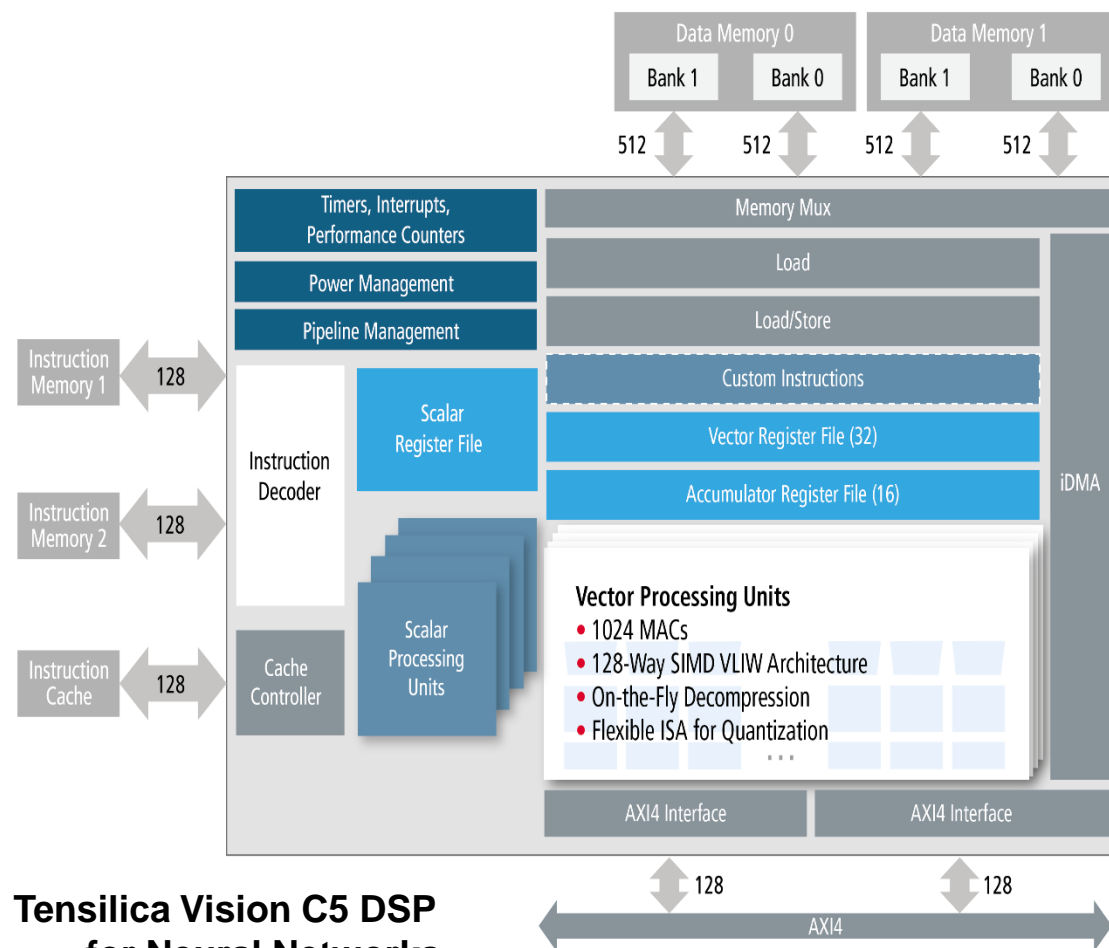
## Cadence Addressing All Market Segments



# Vision C5: DSP Architecture

- Fixed point DSP with 8-bit and 16-bit data type support
- 1024 8x8 MAC or 512 16x16 MAC throughput per cycle
  - Emphasis on high utilization of MACs across range of layer dimensions
- SIMD architecture for high performance vector computing
  - 512-bit vector register file that can work as 1024-bit register (pairing 2 512-bit registers)
    - 128-way SIMD for 8-bit data type, 64-way SIMD for 16-bit data type
  - 1536-bit wide accumulator register file that works as 3072-bit accumulator register (pairing 2 1536-bit accumulator)
- VLIW architecture to exploit instruction level parallelism
  - 88-bit wide VLIW instructions, supports 3 and 4 slot instruction formats
  - Ability to perform load/store, MAC/ALU/SELECT, PACK, decompress operations in parallel
- Dual load/store architecture, capable of two 512-bit loads or one load and one store in parallel
  - Including support for loading unaligned data from memory
  - Special addressing mode for efficient access of 3-D data

# Vision C5: Memory Architecture



**Tensilica Vision C5 DSP  
for Neural Networks**

- TCM based local Data Memory
- Two Data Memory used in ping pong fashion
  - Hide system memory latencies
- Each Data RAM size expected between 128K to 256K depending on application
- Each Data RAM can be banked into 2 or 4 to achieve maximum data access throughput
- Integrated DMA engine to transfer data between local memory and system memory
- Instruction memory supports both caches and TCM
- Two 128-bit AXI interfaces
  - one dedicated to DMA
  - Another one for the processor load/store and instruction memory access to main memory

# Vision C ISA Highlights: Load/Store

- Vector load/store of 64 8-bit elements; can also be viewed as 32 16-bit element load
- Vector loads of 32 8-bit elements, with each element sign/zero extended to 16-bit
- Vector stores of 32 16-bit elements, with each element saturated/truncated to 8-bits
- Vector stores of 32 8-bit elements (low or high half)
- Boolean register load/stores
- Unaligned vector loads of 64 8-bit elements
- Unaligned vector loads of 32 8-bit elements, with element sign/zero extended to 16-bit
- Variable vector load/stores of up to 64 8-bit elements
- Immediate and Index addressing mode, including post increment of address
  - Not all addressing modes supported for all types of loads/stores
- Special loads
  - Load with offset: 8-bit, 16-bit, aligned, unaligned. Example: IVP\_LV2NX8\_ORP
  - Unaligned loads for compressed data.
  - Interleaving loads and stores: Used for MOD vectorization, when data from 2 different locations are to be multiplied with the same set of coefficients. Uses alignment register to help interleave data from two separate loads into a vector register. Example: IVP\_LVINTL16A2NX8\_IP

# Vision C ISA Highlights: ALU, Shifts and Moves

## ALU Support

- 8b, 16b vector add, subtract
- 24b, 48b wide vector add, subtract
- Few variations of Reduction ADD
- Support for addition of upper and lower half of wide accumulator vectors
- Support for different variants of MIN, MAX, CMP (compare) for 8 and 16 bit
- Support for Bitwise AND, OR and XOR for 8 and 16 bit

## Shift Operations

- 8b, 16b arithmetic and logic right shift
- 8b, 16b left shift

## Move Operations

- Vector to vector register move
- Accumulator to accumulator move
- Vector to accumulator and vice versa
- AR to vector register move with replication

# Vision C ISA Highlights: Select Operations

- Extract a scalar 8/16 bit value from vector register to AR
- Arbitrary selection of 8/16-bit elements from two input vector registers to the destination vector register
  - This is the typical SEL operation found in most of our DSPs
- Input can be two vector registers (output is one vector register) or two paired vector registers (output is one paired vector register)
- Special version of the above general operation, using an immediate operand to select from a table of most common select values
- Squeeze and Extract operations illustrated in the example related to avoiding “multiply by zero” data values



# Vision C ISA Highlights: PACK Operations

- PACK operations are used to convert data from the (wider) “accumulator format” to the (narrower) “vector register” format
  - Right shift, saturate, optional rounding
- Vision C supports dual PACKs that pack 2 accumulators to 2 vectors; twice the throughput of VP6
- 256 24-bit elements from accumulator to 256 8-bit elements in vector register
  - Signed and unsigned saturation based on data type of result
- 128 48-bit elements from accumulator to 128 16-bit elements in vector register
- 128 24-bit elements from accumulator to 128 16-bit elements in vector register
- Reduction ADD + PACK: For output of “folded” multiplies
  - 2 way and 4 way reduction ADD before right shift and saturate
  - Result is either 8-bit or 16-bit

# Tensilica® Vision C5 DSP vs NN Accelerator

## Vision C5 DSP

A complete processor that stands on its own:  
**Accelerates all NN layers**

### **Flexible and future-proof solution:**

- Supports variable kernel sizes, depths, input dimensions
- Supports different compression/decompression techniques
- Support for new layers can be added as they evolve

**Main vision/imaging DSP free** to run other applications while NN DSP runs NN

Simple (**single-processing**) programming model for NN

**No need to move** data between NN DSP and main vision/imaging DSP

## NN Accelerator

Built to accelerate **only NN convolution functions**

HW accelerators are mostly designed based on current needs and hence provide a rigid and not future-proof solution

While running NN, **main vision/imaging DSP cannot run other applications**

Complicated **multi-processor** programming model

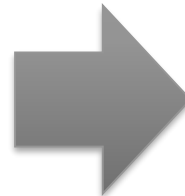
**Need to move** data between NN DSP and main vision/imaging DSP (wastes power)

# Why Vision C5

## Challenge #1: MAC architecture

### Challenge: MAC Architecture

- Efficient implementation of Multiplier accumulator (MAC) architecture
- Achieve high MAC utilization
- MAC architecture has to work across different types of three dimensional convolution
- The MAC architecture has to work with different convolution sizes
  - 11x11, 7x7, 5x5, 3x3, 1x1



### Vision C5 DSP

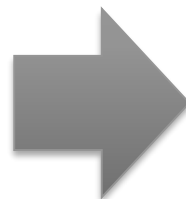
- Specialized dual quad MAC architecture to get almost 100% MAC utilization for inner loops
- Optimizations for very small convolution dimensions
- Enhanced instruction set for multiple vectorization schemes (depth vs width) for best performance in different cases
- Data compression to avoid multiplication by zero

# Why Vision C5

## Challenge #2: Non-Convolution Layers

### Challenge: Multiple Layers

- Enhancements for high performance processing of non-convolution layers
- Every convolution and FC layer followed by non convolution layer
- Bit dept flexibility per layer:  
Convolution layer vs Normalization



### Vision C5 DSP

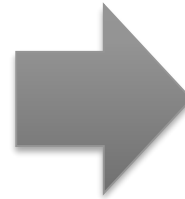
- Specific set of ALU operations for enhanced non-convolution layers
- Fusion of multiple layers for higher performance overall with specialized instruction set
- Mixed 8bit and 16bit precision

# Why Vision C5

## Challenge #3: Quantization

### Challenge: Quantization

- Fixed point support needed for 8-bit to 16-bit both
- Requires quantization support using networks trained in floating point deployed for inference in fixed point
- Maintain high accuracy after quantization



### Vision C5 DSP

- 1024 8-bit and 512 16-bit MAC
- On the fly precision mixing, Eg: 8bit for conv gives higher performance and 16bit for normalization provides better accuracy
- Enhanced operations for efficient quantization

# Why Vision C5

## Challenge #4: Memory Bandwidth

### Challenge: Memory Usage

- Memory bandwidth is limited and often shared
- Access to memories must be efficient and necessary
- Reduce data manipulation impact on processor cycles to keep processing efficiency high

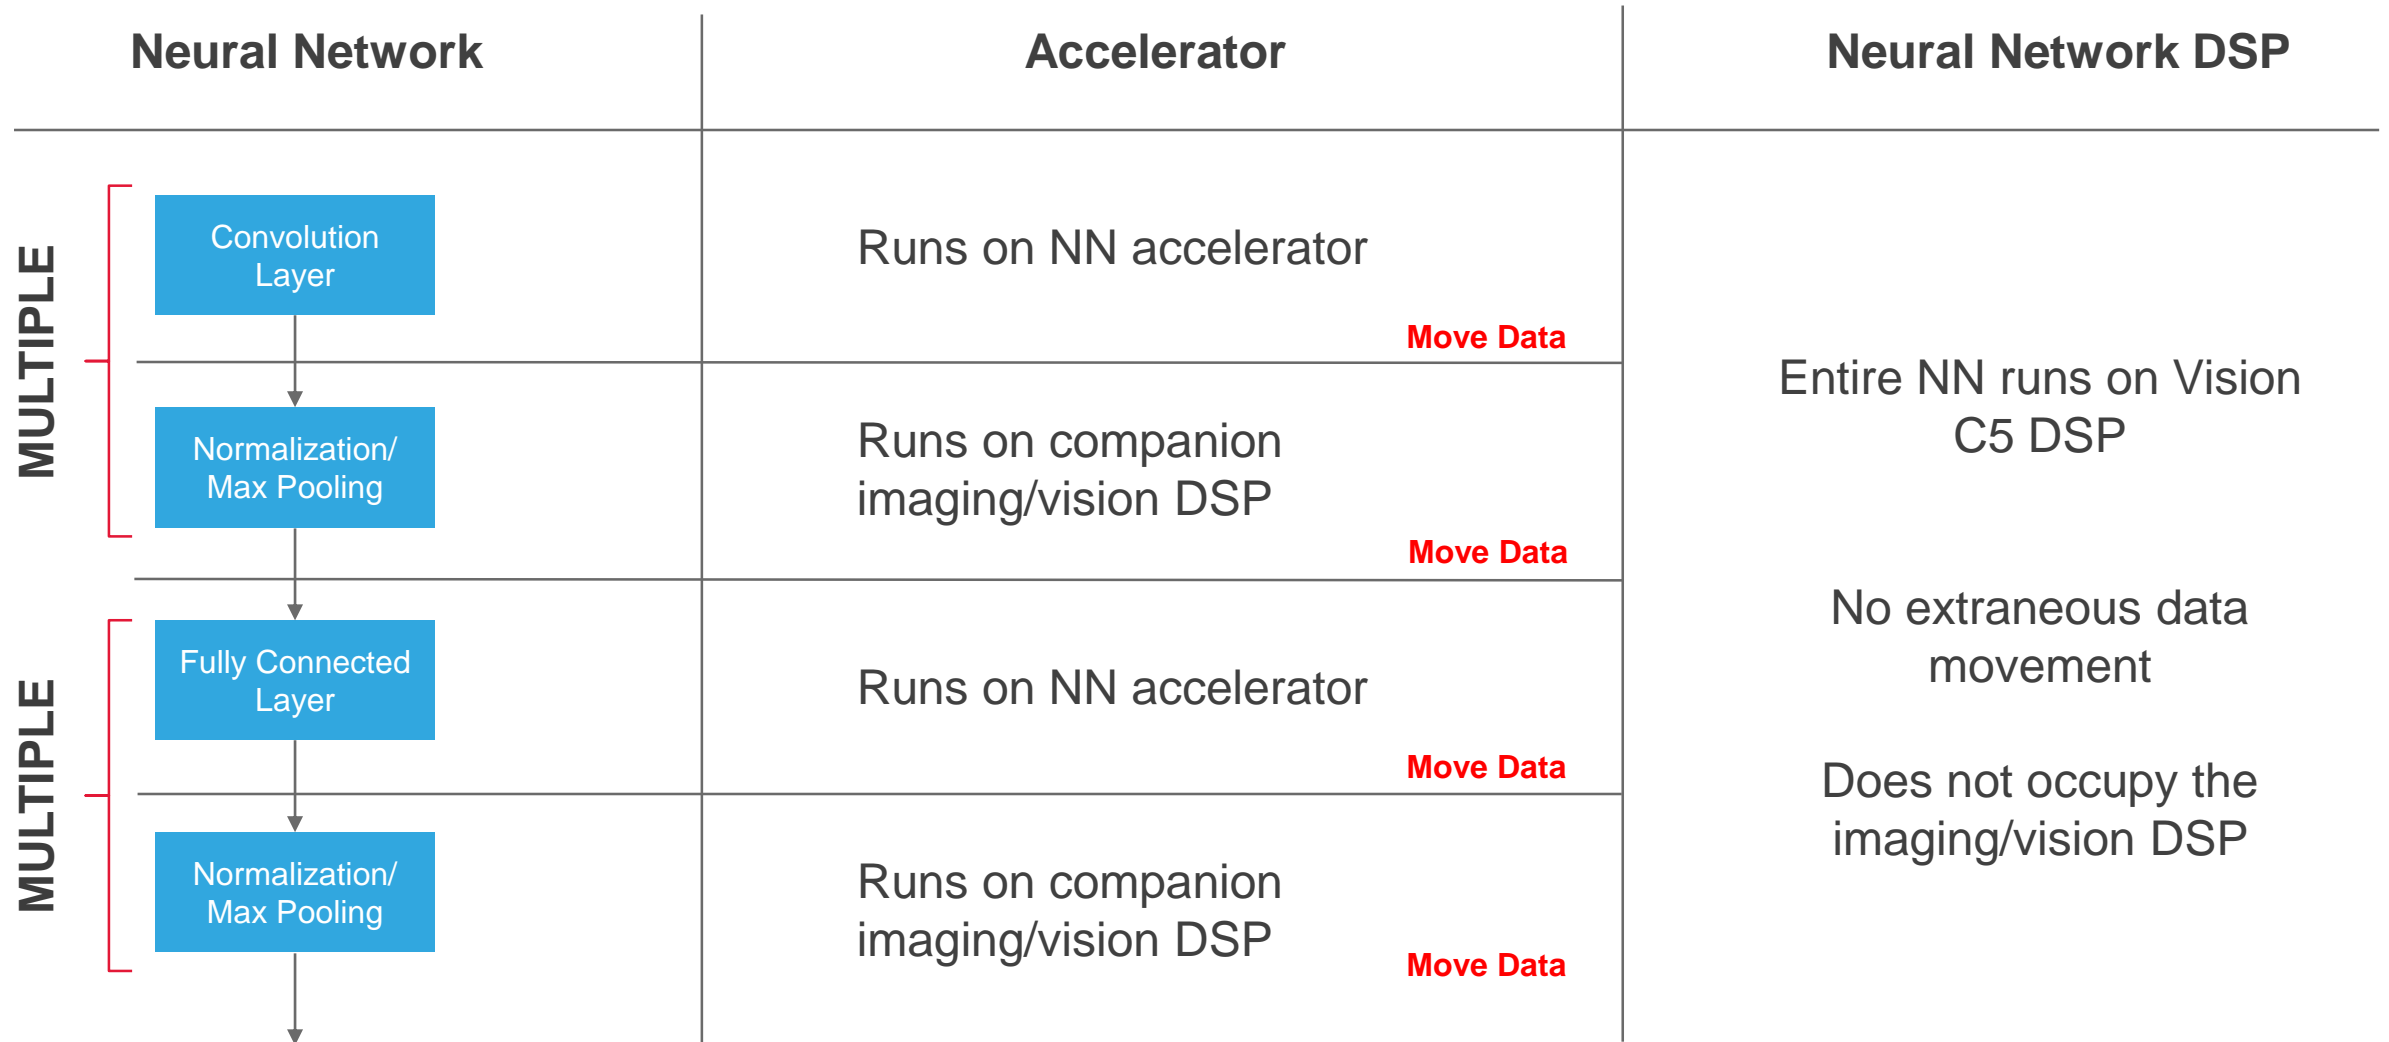


### Vision C5 DSP

- Integrated DMA:  
No need for controller, can manage its own data movement
- Ping-pong memory buffering hides DMA latency
- Rich and wide set of select instructions for data manipulation
- On the fly weight decompression for bandwidth reduction

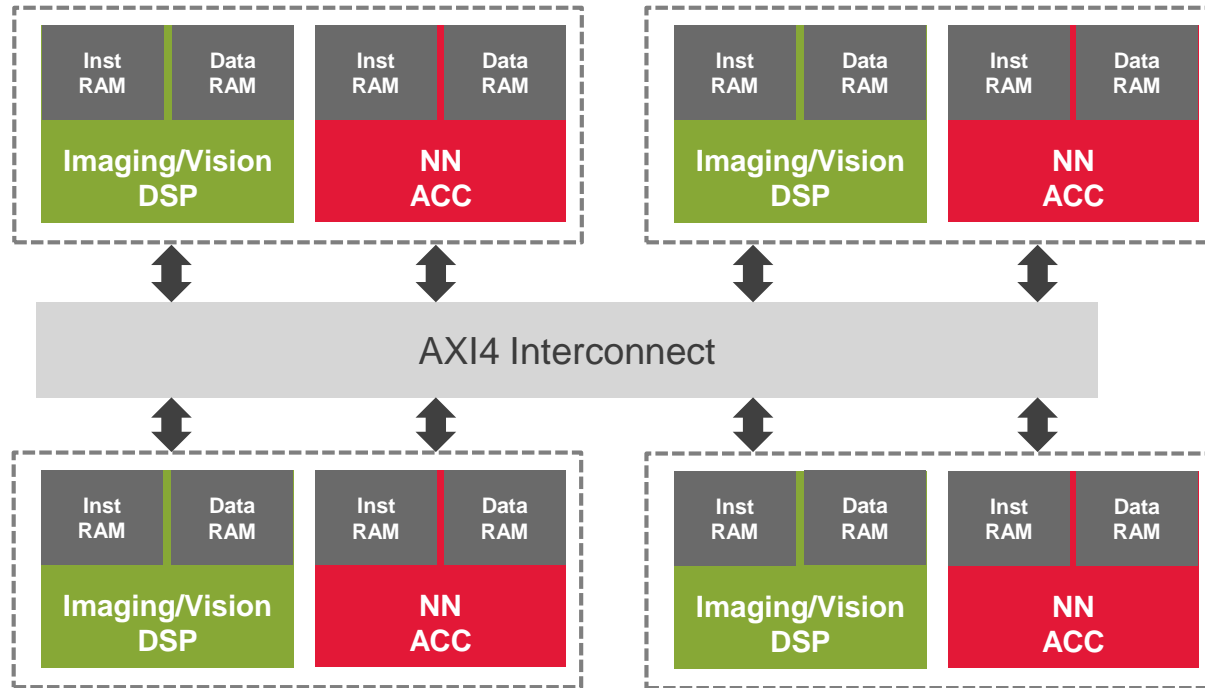


# Tensilica® Vision C5 DSP vs NN Accelerator

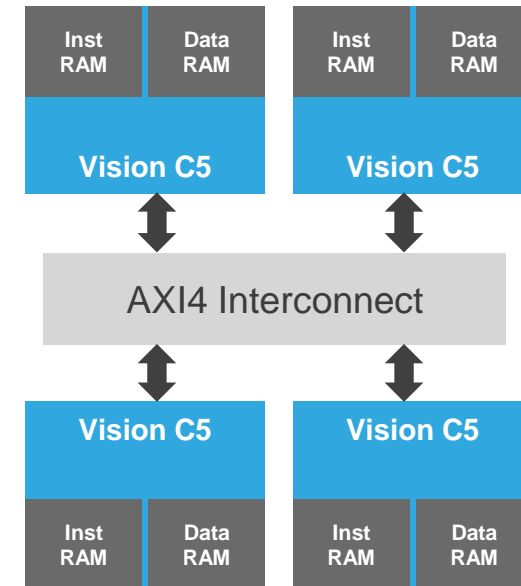


# Multi-Core Solution with Tensilica® Vision C5 DSP vs NN Accelerator

## Multi-Core with Vision DSP + NN Accelerator



## Multi-Core with Vision C5 DSP



- Vision C5 DSP scales elegantly
- NN accelerator approach requires imaging/vision DSP with each core; increased area and power

# Vision C5 Area: TSMC 16FFC 9-Track

## Post layout Core Cell, Memory Cell and final Floor Plan

| Metric                   | 500MHz                                   | Notes   |
|--------------------------|--|---|
| Vision C5 reference with | 64KB 4-way I\$<br>2 x 128 KB Data RAMs   |   |
|                          | RVT only –<br>Vdd 0.8 V                  |   |
| Memory compiler:         | ts1n16ffcllsvta64x20m4sw<br>version 100a |   |
| Standard Cell Core Area  | <b>0.895 mm<sup>2</sup></b>              | Post layout standard cell core area of Vision C5 not including utilization ratio (including iDMA) |
| Memory Area              | 0.334 mm <sup>2</sup>                    | Area depends on amount of memory, memory compiler   |

Vision C5 8b MAC/mm2 = 2.8x Vision P6 8b MAC/mm2  
 Vision C5 16b MAC/mm2 = 5.6x Vision P6 16b MAC/mm2

# Vision C5 Area: TSMC 28HPC+ 9-Track

## Post layout Core Cell, Memory Cell and final Floor Plan

| Metric                   | 500MHz                                 | 750MHz                                 |  |
|--------------------------|--|--|--|
| Vision C5 reference with | 64KB 4-way I\$<br>2 x 128 KB Data RAMs | 64KB 4-way I\$<br>2 x 128 KB Data RAMs |  |
|                          | RVT only –<br>Vdd 0.9 V                | RVT only –<br>Vdd 0.9 V                |  |
| Memory compiler:         | Tcbn28hpcplusbwp35p140<br>version 110c | Tcbn28hpcplusbwp35p140<br>version 110c |  |
| Standard Cell Area       | <b>1.99 mm<sup>2</sup></b>             | <b>2.679 mm<sup>2</sup></b>            | Post layout standard cell area of<br>Vision C5<br>(including iDMA) |
| Memory Area              | 0.571 mm <sup>2</sup>                  | 0.571 mm <sup>2</sup>                  | Area depends on amount of memory,<br>memory compiler               |

Vision C5 8b MAC/mm2 = 2.8x Vision P6 8b MAC/mm2  
 Vision C5 16b MAC/mm2 = 5.6x Vision P6 16b MAC/mm2

# AlexNet Performance Estimate on Vision C5 (8bit Data and 8 bit Weights)

| AlexNet Details  |           |
|------------------|-----------|
| Input ROI        | 227x227x3 |
| Number of Layers | 14        |

| Vision P6 FPS/GHz |             | Vision C5 FPS/GHz |             |
|-------------------|-------------|-------------------|-------------|
| FPS               | DDR Latency | FPS               | DDR Latency |
| 214               | 100 cycles  | 700               | 100 cycles  |

| VC5 AlexNet Weighted Power (mW/MHz) |      |
|-------------------------------------|------|
| TSMC 16FFC, RVT, 0.8V               | 0.65 |
| TSMC 28HPC+, RVT, 0.9V              | 1.05 |

- **Performance:** Vision C5 = **3.27x** Vision P6
- **Efficiency (Perf/mm2):** Vision C5 = **2.3x** Vision P6
- **Efficiency (Perf/mW):** Vision C5 = **1.5x** Vision P6
- FPS numbers are from manual optimized implementation
- Batching is used in layers that are potentially memory limited
- Performance results are preliminary pending further optimizations

# ResNet50 Performance Estimate on Vision C5

(8bit Data and 8 bit Weights)

| ResNet50 Details |           |
|------------------|-----------|
| Input ROI        | 224x224x3 |
| Number of Layers | 223       |

| Vision P6 FPS/GHz |             | Vision C5 FPS/GHz |             | VC5 ResNet50 Weighted Power (mW/MHz) |      |
|-------------------|-------------|-------------------|-------------|--------------------------------------|------|
| FPS               | DDR Latency | FPS               | DDR Latency | TSMC 16FFC, RVT, 0.8V                | 0.82 |
| 33                | 100 cycles  | 129               | 100 cycles  | TSMC 28HPC+, RVT, 0.9V               | 1.31 |

- **Performance:** Vision C5 = **3.7x** Vision P6
- **Efficiency (Perf/mm2):** Vision C5 = **2.6x** Vision P6
- Vision P6 is measured from XNNC implementation and Vision C5 is estimated from hand written kernels extrapolated to all layers
- Performance results are preliminary pending further optimizations
- Reference memory configuration of 64KB 4-way I\$ and 2x128 Data RAM has been used
- Batching is used in any layer that is potentially memory limited



# InceptionV3 Performance Estimate on Vision P6 and Vision C5 (8bit Data and 8 bit Weights)

| Inception V3 Details |           |
|----------------------|-----------|
| Input ROI            | 299x299x3 |
| Number of Layers     | 110       |

| Vision P6 FPS/GHz |             | Vision C5 FPS/GHz |             |
|-------------------|-------------|-------------------|-------------|
| FPS               | DDR Latency | FPS               | DDR Latency |
| 30                | 100 cycles  | 103               | 100 cycles  |

- **Performance:** Vision C5 = **3.4x** Vision P6
- **Efficiency (Perf/mm2):** Vision C5 = **2.4x** Vision P6
- Vision P6 is measured from XNNC implementation and Vision C5 is estimated from hand written kernels extrapolated to all layers
- Reference memory configuration of 64KB 4-way I\$ and 2x128 Data RAM has been used
- C5 results are preliminary pending further optimizations

# Vision C5 Benchmark (Early Estimates)

| Layer              | Dimensions           |                   |                   | Precision | Vector Approach | MAC Utilization (in %)          |                            |
|--------------------|----------------------|-------------------|-------------------|-----------|-----------------|---------------------------------|----------------------------|
|                    | Input<br>(W x H x D) | Number of Kernels | Kernel<br>(W x H) |           |                 | Inner loop<br>(static schedule) | Kernel<br>(with mem model) |
| Custom Convolution | 960x5x32             | 16                | 5x5               | 8b x 8b   | MOW             | 100                             | 86                         |
| Custom Convolution | 480x5x32             | 16                | 5x5               | 16b x 16b | MOW             | 100                             | 86                         |
| Custom Convolution | 1020x5x32            | 16                | 5x5               | 8b x 8b   | MOW             | 100                             | 93                         |
| Custom Convolution | 160x6x64             | 64                | 5x5               | 8b x 8b   | MOW             | 100                             | 71                         |
| Custom Convolution | 160x6x64             | 64                | 5x5               | 8b x 8b   | MOD             | 89                              | 74                         |
| Res2a_branc2b      | 56x16x64             | 64                | 3x3               | 8b x 8b   | MOW             | 100                             | 75                         |
| Res2a_branc2b      | 56x16x64             | 64                | 3x3               | 8b x 8b   | MOD             | 89                              | 75                         |
| Resnet 1x1         | 56x56x256            | 8                 | 1x1               | 8b x 8b   | MOD             | 94                              | 75                         |
| Resnet 1x1         | 28x28x256            | 8                 | 1x1               | 16b x 16b | MOD             | 89                              | 73                         |
| FC (32 Batch)      | 2x1x256              | 128               | 2x1x256           | 8b x 8b   | MOD             | 94                              | 66                         |

- Tile sizes are based on 2 x 128K Data RAM
- Performance is for tiles in local memory
- DMA is assumed to be in background and not included in above cycles

# Vision C5 DSP vs Nvidia TX1

AlexNet

Up to 6X\* faster

Inception V3

Up to 9X\*\* faster

ResNet50

Up to 4.5x\*\*\* faster

Note:

<https://arxiv.org/pdf/1605.07678v2.pdf>

Both cores running at 690MHz on 16nm

\* AlexNet data with 8 batch

\*\* Inception V3 data with single batch

\*\*\* ResNet50 with batching

# Vision P6 vs Vision C5

|  | Vision P6  | Vision C5   |
|--|--|---|
| Focus  | Imaging and NN                                     | NN  |
| MAC (8x8)  | 256  | 1024  |
| MAC (16x16)  | 64   | 512   |
| Single and Half Precision VFPU<br>(optional)       | 32 way FP16<br>16 way FP32                         | Not Required for Inference  |
| Accumulators<br>(to support higher MAC capability) | 4 x 1536b  | 8 x 3072b   |
| MAX SIMD Width                                     | 64 way SIMD  | 128 way SIMD  |
| Special Features                                   | Scatter Gather<br>(needed by Imaging Applications) | <ul style="list-style-type: none"> <li>• On the Fly Decompression Support</li> <li>• Special addressing modes</li> <li>• Richer set of convolution multipliers (signed and unsigned)</li> <li>• Extensive data rearrangement and selection</li> </ul> |

# Cadence Tensilica Vision C5 DSP Summary

For all neural network inference applications

## DSP optimized to process all neural network inference layers

- Instruction set and DSP architecture designed specifically for NN
- Not just a “convolution accelerator”

## Scalable from 1 TeraMAC/sec computational capacity in less than 1mm<sup>2</sup>

- Easily scales to multi-TMAC/sec systems, no limit to the number of Vision C5 DSPs
- Designed for high-availability (always-on) neural network computational needs

## Flexible and programmable

- To meet evolving neural network requirements

## Targeted at surveillance, automotive, drone and mobile/wearable markets

- Optimized for vision, radar/lidar and fused-sensor applications



# Development Tools



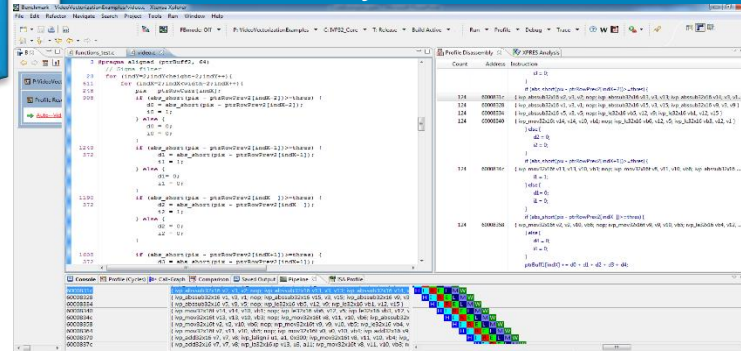
# Vision Family Development Tools

## Developed on rich Xtensa heritage

- >17 year history of Xtensa architecture
- Billions of core in production

## Seamless development environment

- Eclipsed based IDE GUI
- Xtensa C/C++ (XCC) Compiler with auto-vectorization
- GNU Software Toolkit (Assembler, Linker, Debugger, Profiler)
- Vision Library

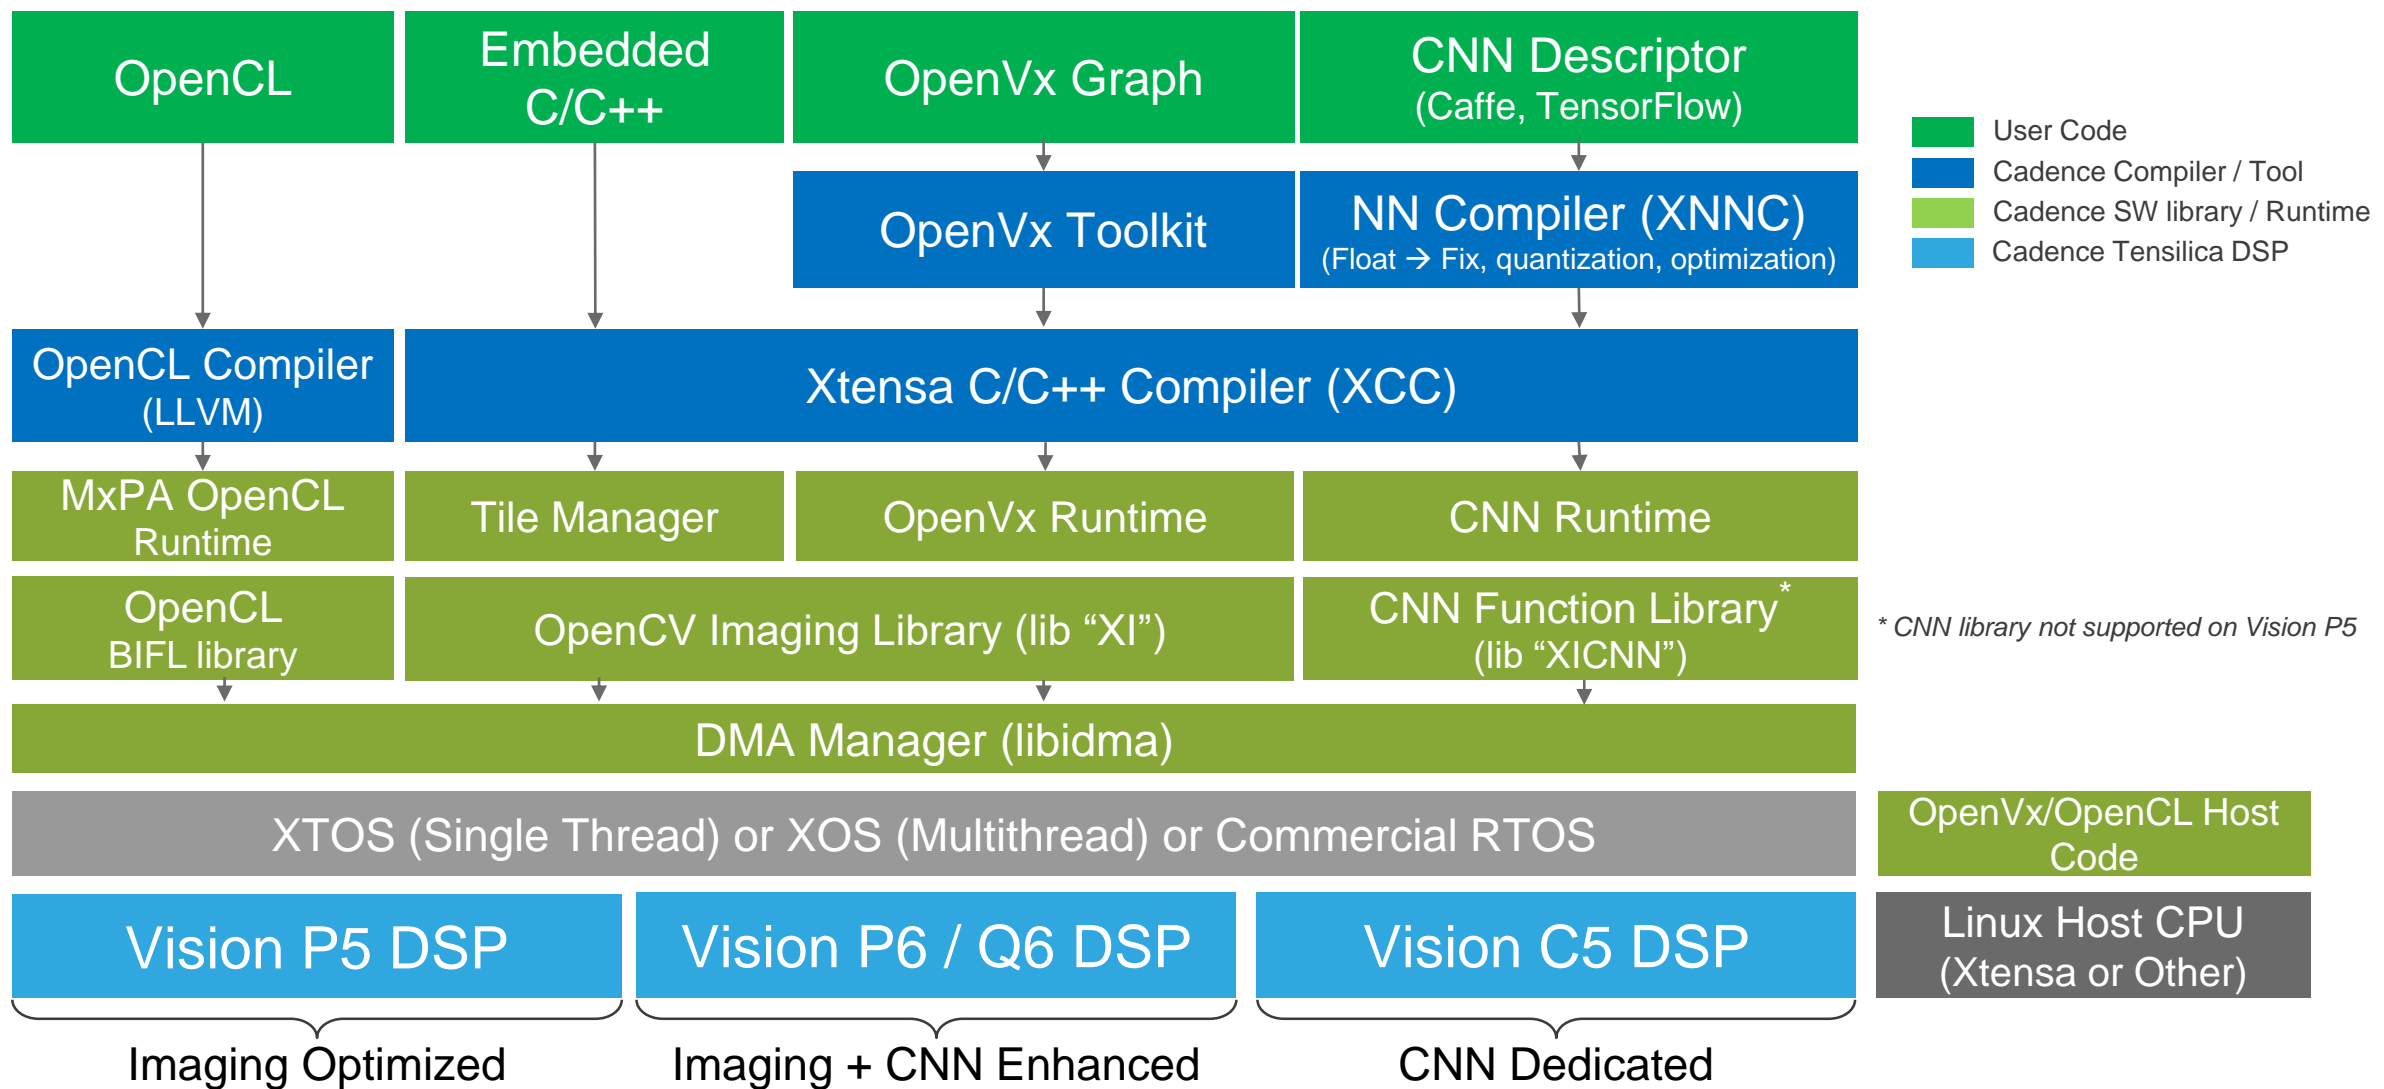


## Complete package for integration

- RTL, EDA script
- Cycle accurate Instruction Set model
- Fast Function Simulator (TurboXIM)
- XTSC System C system modeling

# Tensilica: Comprehensive Vision Software & Hardware Solutions

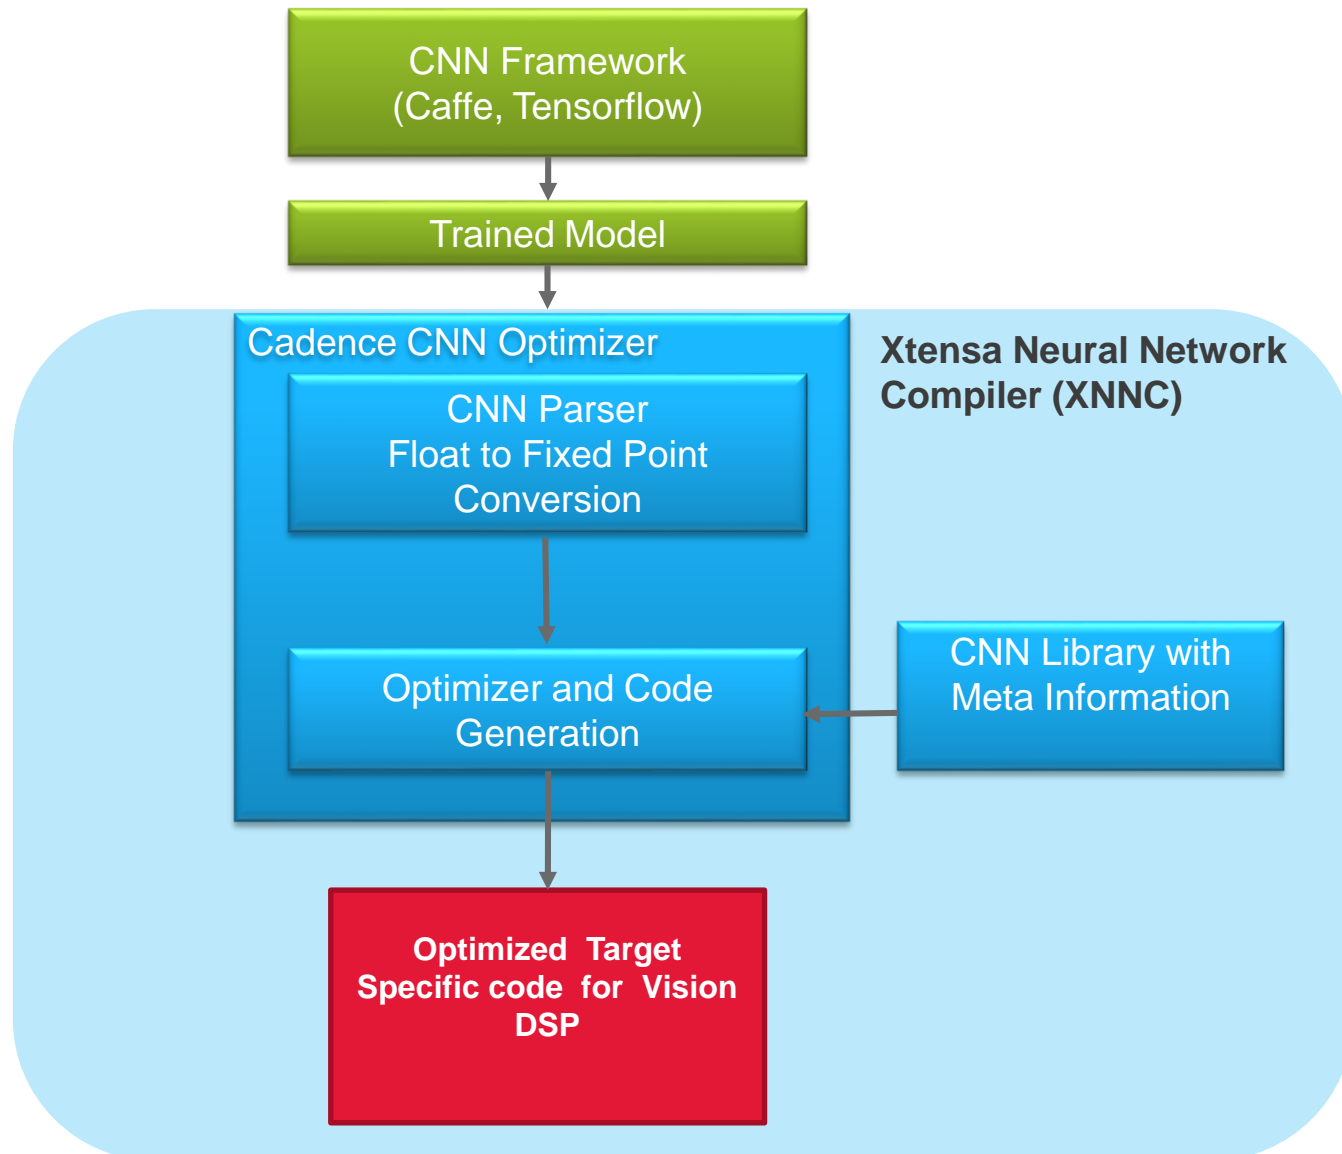
*Full Ecosystem of Software Frameworks and Compilers for all Vision and Imaging Programming Styles*





# Xtensa Neural Network Compiler (XNNC)

(Starting From Vision P6)

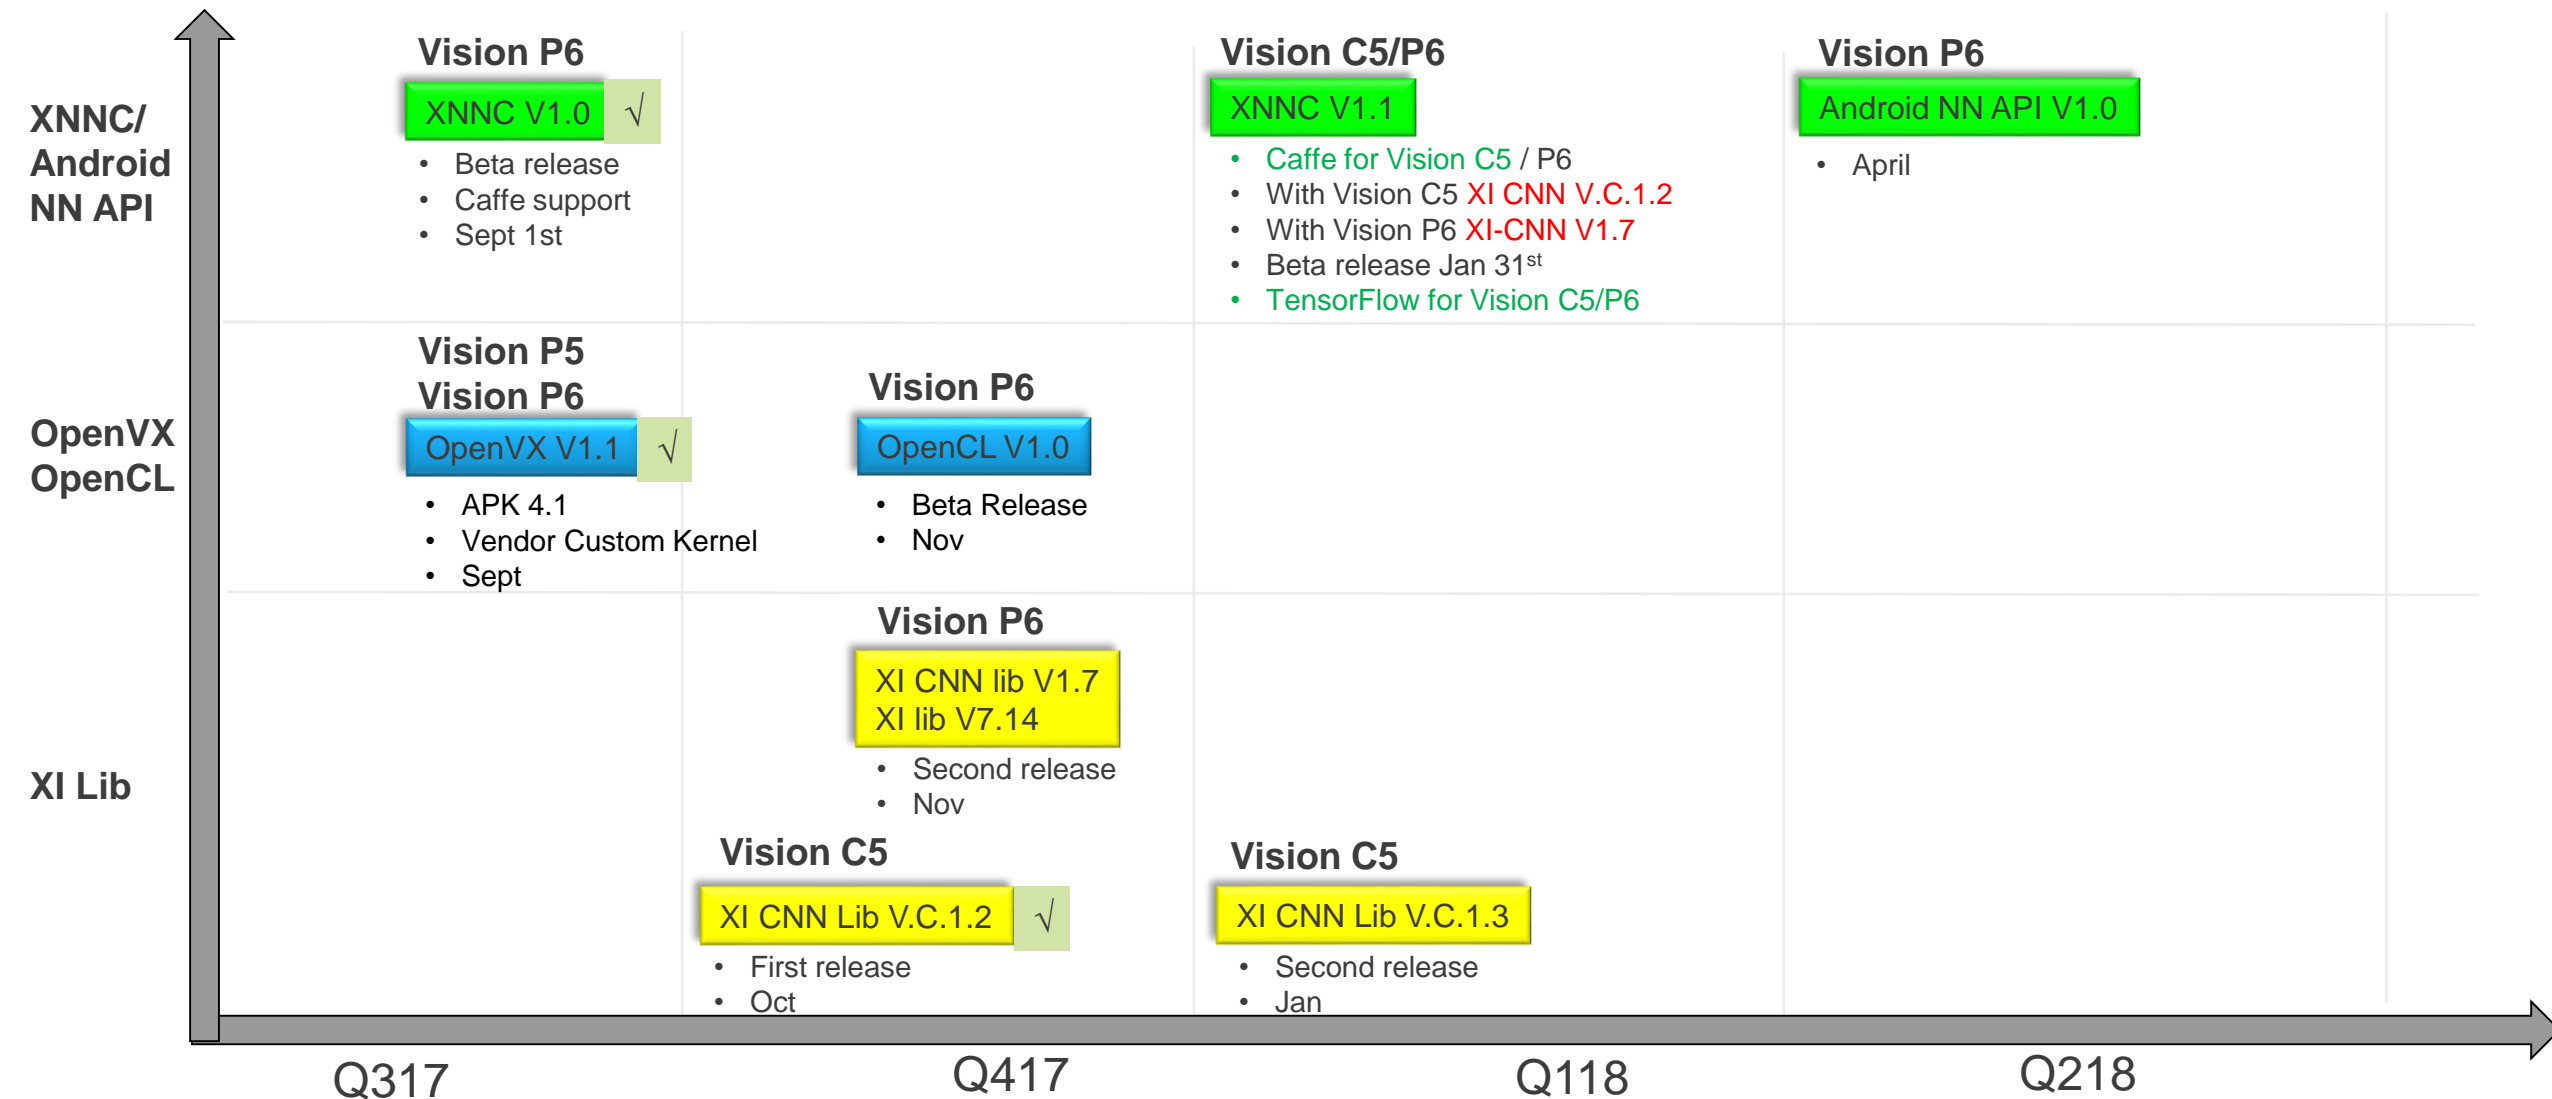


- Connects to existing industry CNN frameworks by using their Trained Model descriptions
- and **auto-generates Trained Model optimized code for Cadence CNN DSPs**

- Three Major components to XNNC
- CNN Parser: Float to Fixed Point conversion
- CNN Code Generation and Optimization
- CNN Library for Vision DSP

- First CNN Framework support: Caffe, followed by Tensorflow
- For both Vision P6 and Vision C5

# Vision DSP SW Roadmap



# Applications Across Market and Cadence Ecosystem

|                             | Mobile | Auto | Security | Gaming | Wearable | Partners (Sources)   |  |
|-----------------------------|--------|------|----------|--------|----------|--|--|
| Face Detect                 |        |      |          |        |          | Irida Labs, Cadence  | Developed by Cadence<br>Irida Labs                         |
| People Detection (HOG)      |        |      |          |        |          | Uurmi System (Now Mathworks),<br>MultiCoreware, Cadence (demo) | Developed by Cadence                                       |
| DoG Differences of Gaussian |        |      |          |        |          | BDTi, Cadence (XI Lib)   | Parameterized and Optimized Version                        |
| CNN                         |        |      |          |        |          | Multicoreware, BDTi  | CNN prediction done by Multicoreware<br>and BDTi<br>(Yolo) |
| Lane Departure Warning      |        |      |          |        |          | Cadence (Demo)   | Available  |
| Fog removal                 |        |      |          |        |          | Uurmi System (Now Mathworks)                                   | Available on Vision P5 (Uurmi)                             |
| Traffic Sign Detection      |        |      |          |        |          | Cadence (Demo)   | Developed by Cadence (CNN)                                 |
| Stereo Sensors              |        |      |          |        |          | Uurmi(Now Mathworks), Multicoreware                            | Available on Vision P5 (Uurmi)                             |
| Super Resolution            |        |      |          |        |          | Alamalence, Uurmi Systems(Now<br>Mathworks)                    | Available on Vision P5                                     |
| Video Stabilizer            |        |      |          |        |          | Irida Labs, Morpho, Uurmi<br>Systems(Now Mathworks)            | Morpho   |
| Video WDR                   |        |      |          |        |          | Morpho   | Morpho   |
| Low-Light                   |        |      |          |        |          | Irida labs, Morpho   | Available (Irida labs)                                     |
| 3D Noise filtering          |        |      |          |        |          | Morpho, Irida labs   | Available (Irida labs)                                     |
| Face & Voice Authentication |        |      |          |        |          | Sensory  | Available  |
| 360 Defish                  |        |      |          |        |          | Multicoreware  | Available  |

# New Engagements

- Vangogh Imaging
  - For SLAM Algorithm
- Arcsoft
  - Face detection and face beautification
- Corephotonics
  - Stereo Image fusion
- ReadSense
  - Face recognition/People recognition
- Thundersoft
  - Computer Vision and Android

# Summary

## Imaging DSP and Neural Network DSP

- Market needs both Imaging DSP and Neural Network DSP
- Imaging processing requirement continues to increase: higher resolution and dual sensor
- Neural Network are evolving – computational capacity continue to increase

## Vision C5 DSP for Neural Network

- Complete, standalone DSP that runs all layers of NN : Not an accelerator
- General purpose and programmable
- 1 TMAC/Sec computational capacity: 4X MAC capacity compare to Vision P6

## Vision P6 DSP

- Up to 4X neural network performance compare to Vision P5
- 4X MAC capacity
- Easy migration of GPU floating-point code, with optional 16-bit floating-point unit
- Up to 4X performance improvement on well-known imaging/vision benchmarks

Relative numbers presented are in comparison with Tensilica Vision P5 DSP

**cā dence<sup>®</sup>**