

2016

www.ceva-dsp.com

Agenda



- Glossary
- Median Filter 3X3
 - Algorithm Analysis
 - Efficient VEC-C Code
 - Performance
- Gaussian3X3
 - Gaussian3X3 Kernel Overview
 - Gaussian3X3 Vector Operators Code
 - Gaussian3X3 load Operation Optimizations
 - Gaussian3X3 Sliding window code

Agenda



Glossary

- Median Filter 3X3
 - Algorithm Analysis
 - Efficient VEC-C Code
 - Performance

Gaussian3X3

- Gaussian3X3 Kernel Overview
- Gaussian3X3 Vector Operators Code
- Gaussian3X3 load Operation Optimizations
- Gaussian3X3 Sliding window code

Glossary



- Vector types newly defined CEVA vector native types to be used in vector operation and mapped to the VPU vectors resources
 - 8 integer vector

	8i.Xv											
	i3 i2					i1			i0			
127		96	95		64	63		32	31		D	
	i7		i6		i5			i4				
255		224	223		192	191		160	159		128	

16 short vector

	vX.s16																						
	s7			s6			s5			s4			s3			s2			s1			s0	
127		112	111		96	95		80	79		64	63		48	47		32	31		16	15		0
	s15			s14			s13			s12			s11			s10			s9			s8	
255		240	239		224	223		208	207		192	191		176	175		160	159		144	143		128

32 byte vector

							vX.	c32							
c15	c14	c13	c12	c11	c10	c9	c8	c7	c6	c5	c4	c3	c2	c1	c0
127 12	119 112	111104	103 95	95 88	87 80	79 72	71 64	63 55	55 48	47 40	39 32	31 24	23 16	15 8	7 0
c31	c30	c29	c28	c27	c26	c25	c24	c23	c22	c21	c20	c19	c18	c17	c16
25524	8 247 240	239232	231224	223216	215203	207200	199 192	191 184	183 176	175 163	167 160	159 152	151 144	143136	135 128

Native 256-bit vectors							
char32	uchar32						
short16	ushort16						
int8 uint8							
128-bit vectors							
char16	uchar16						
short8	ushort8						
512-bit vectors							
short32	ushort32						
int16	uint16						

Glossary



Vector operators – Simple way of coding an optimized Vec-C code, using vector operators:

```
uchar32 v00, v01, v02;
uchar* p_in_u8 = (uchar *)&p_u8Src[i * stride];
v00 = *(uchar32*)p_in_u8;
p_in_u8++;
v01 = *(uchar32*)p_in_u8;
p_in_u8++;
vacc = (ushort32)v00 * 2 + (ushort32)v01 * 4;
```

- ► The code is written in a plain way
- Vector operators are used with the vector type as done for other native C language types [char, short or int].
- Ensures fast and simple code development
- Does not give access to all instruction set features

Glossary



▶ Vec-C – vector intrinsic that have one to one mapping with the instruction set:

```
uchar32 v00, v01, v02;
uchar* p_in_u8 = (uchar *)&p_u8Src[i * 32];
uint16 vacc0;
v00 = *(uchar32*)p_in_u8;
p_in_u8+=s32SrcStep;
v10 = *(uchar32*)p_in_u8;
p_in_u8+=s32SrcStep;
vacc0 = (uint16)vswmpy5(v00, v00, v_coeff, config0);
```

- Enables to use CEVA-XM4 rich instruction set features from a C level environment.
- Provides superb performance per invested coding time.

Agenda



Glossary

♦ Median Filter 3X3

- Algorithm Analysis
- Efficient VEC-C Code
- Performance

Gaussian3X3

- Gaussian3X3 Kernel Overview
- Gaussian3X3 Vector Operators Code
- Gaussian3X3 load Operation Optimizations
- Gaussian3X3 Sliding window code

Median3X3



- Algorithm Analysis
- Efficient VEC-C Code
- ***Performance**

Median3*3 Introduction



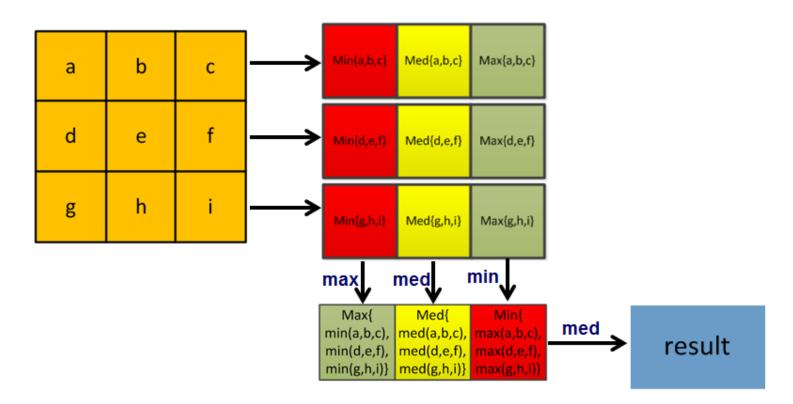
Dst[x,y] = Median{

$$Src[x-1,y]$$
 , $Src[x,y]$, $Src[x+1,y]$

$$Src[x-1,y+1], Src[x,y+1], Src[x+1,y+1]$$

Implementation





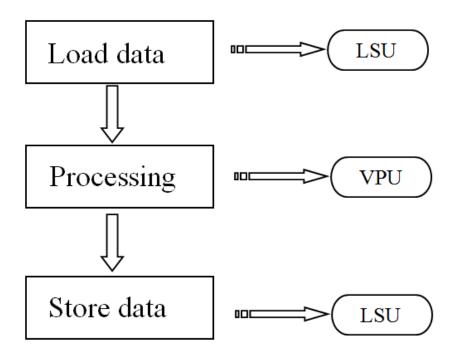
Median3X3



- Algorithm Analysis
- Efficient VEC-C Code
- **❖Performance Analysis**

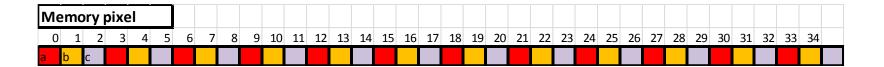
Data Process Flow

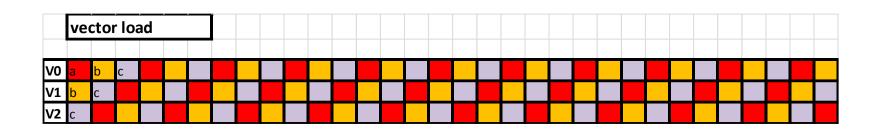




Loading Data







vldov((uchar32*)p_in_u8, v0, v1, v2, vdummy);

Processing and Store



```
v0min = vmin(v1, v2, v3);
v0med = vmed(v1, v2, v3);
v0max = vmax(v1, v2, v3);
v1min = vmin(v4, v5, v6);
v1med = vmed(v4, v5, v6);
v1max = vmax(v4, v5, v6);
v2min = vmin(v7, v8, v9);
v2med = vmed(v7, v8, v9);
v2max = vmax(v7, v8, v9);
v1 = vmax(v0min, v1min, v2min);
v2 = vmed(v0med, v1med, v2med);
v3 = vmin(v0max, v1max, v2max);
uchar32 result = vmed(v1, v2, v3);
vst(result, p dst, vprMask2);
```

Performance



	median3x3
load instruction	vldov(p_src, v1, v2, v3, vdummy)
vpu instruction	vmax(v1, v2, v3)
	vmed(v1, v2, v3)
	vmin(v1, v2, v3)
store instruction	vst(result, p_dst, vprMask)

Agenda



- Glossary
- Median Filter 3X3
 - Algorithm Analysis
 - Efficient VEC-C Code
 - Performance

⋄Gaussian3X3

- Gaussian3X3 Kernel Overview
- Gaussian3X3 Vector Operators Code
- Gaussian3X3 load Operation Optimizations
- Gaussian3X3 Sliding window code

Gaussian3X3 Kernel Overview



Gaussian3x3 coefficients {2,4,2,4,8,4,2,4,2}

$$\triangleright \mathsf{Res} = \frac{1}{32} \sum_{i=0}^{i < 3} \sum_{j=0}^{j < 3} Source_{pixel(Xpos+i,Ypos+j)} * Guas_coef(i,j)$$

2	4	2
4	8	4
2	4	2

C reference code to be optimized

```
for (i = 0; i < u32M; i++)
     uint acc = 0:
     acc += p in u8[-1 - s32SrcStirde] * s8Kernel[0];
     acc += p in u8[0 - s32SrcStirde] * s8Kernel[1];
     acc += p_in_u8[1 - s32SrcStirde] * s8Kernel[2];
     acc += p in u8[-1 ] * s8Kernel[3];
     acc += p_in_u8[0 ] * s8Kernel[4];
     acc += p in u8[1 ] * s8Kernel[5];
     acc += p in u8[-1 + s32SrcStirde] * s8Kernel[6];
     acc += p in u8[0 + s32SrcStirde] * s8Kernel[7];
     acc += p_in_u8[1 + s32SrcStirde] * s8Kernel[8];
     res = (acc >> ps) &0xFFFF;
     sat = (res < 0) ? 0 : (res>255) ? 255 : (uchar)res;
     p_out_u8[i] = sat:
     p_in_u8++;
```

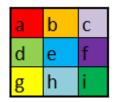
Gaussian3X3 Vector Operators Code



- ► Check that CEVA_OPPERATORS is open in the gaussian3x3_lib.h file.
- Check that the following defines are closed:
 - CEVA_VLDOV and CEVA_SLIDING_WINDOW
- Algorithm flow:

```
Load operation:
                                                         v00 = *(uchar32*)p_in_u8;
                                                         p_in_u8++;
    Vmac3 operation:
                                                         vacc = (ushort32)v00 * 2 + (ushort32)v01 * 4;
                                                         vout = (uchar32)((vacc + (ushort32)v22 * 2 + (ushort32)v22 * 0) >> 5);
    Store operation:
                                                         *(uchar32*)p out u8 = vout;
[slice = 32 Bytes vector]
For (.. ) Horizontal loop {
         calculate input pointer of source slice
         calculate output pointer of destination slice
          For (.. ) Vertical loop
                      read 9 input vectors
                      calculate the filter (including scaling)
                      store results
```

Pixels Data





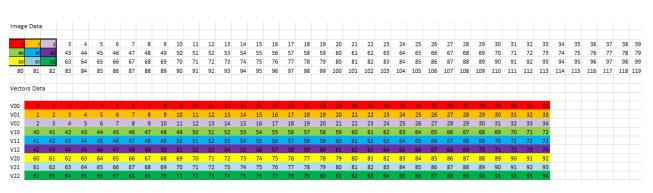


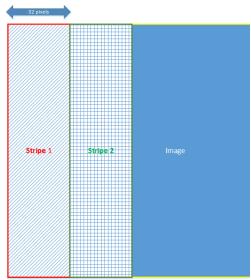
Gaussian3X3 Vector Operators Code



Code Flow

- ▶ The Kernel will work on the source buffer (size [M rows] x [N cols])
 - → In this example size = 128X64
- Scalar code process 9 pixels → Vectorized code will process 9 vectors of 32 Bytes each.
- ▶ Outer loop will scan the buffer on the horizontal way (On cols). Inner loop on the Vertical (On rows)







- Loading the data
 - ► Each load instruction loads 32 bytes into one vector uchar32
 - ► Each calculation produces 32 output points
 - ► The code flow performs 9 vector loads
 - Vector load example (3 loads out of 9)

```
uchar32 v00, v01, v02;
v00 = *(uchar32*)p_in_u8;
p_in_u8++;
v01 = *(uchar32*)p_in_u8;
p_in_u8++;
v02 = *(uchar32*)p_in_u8;
p_in_u8 += s32SrcStrideM2;
```

V00	V01	V02
V10	V11	V12
V20	V21	V22

- ▶ Note the different pointer post modification:
 - Advance to the next pixel
 - Advance to the next pixel line



- ► Gaussian calculation —filter implementation
 - We will use vector operators to perform the filter calculation
 - ▶ Vout = (V00 * 2 + V01 * 4 + V02 * 2 + V10 * 4 + V11 * 8 + V12 * 4 + V20 * 2 + V21 * 4 + V22 * 2 + V22 * 0) >> 5;

```
uchar32 v00, v01, v02;
uchar32 v10, v11, v12;
uchar32 v20, v21, v22;
Uchar32 vout:
ushort32 vacc;
vacc = V00 * 2 + V01 * 4:
vacc += V02 * 2 + V10 * 4;
vacc += V11 * 8 + V12 * 4:
vacc += V20* 2 + V21 * 4;
vout = ((vacc + V22 * 2 + V22 * 0) >> 5);
```

2	4	2
4	8	4
2	4	2

V00	V01	V02
V10	V11	V12
V20	V21	V22



- Gaussian Calculation
 - ► Each Gaussian calculation contains 9 multiplications
 - By utilizing VMAC3 instructions 32 results can be achieved in 5 cycles
 - Calculation using multiple accumulate implementation code
 uchar32 V00, V01, V10, V11;
 ushort32 vacc;
 vacc = (ushort32)V00* 2 + (ushort32)V01 * 4; //Line1 → translated to vmad
 vacc += (ushort32)V10 * 2 + (ushort32)V11 * 4;//Line2 → translated to vmac3



```
VPU1.vmad {init} v23.uc32, r6.uc, v9.uc32, r4.uc, #0x0, v32.us16, v33.us16
VPU1.vmac3 v21.uc32, r6.uc, v22.uc32, r4.uc, v32.s16, v33.s16
```

- ▶ What is the difference between line1 and line2, which operation line2 substitute?
- Why vacc variable is defined as ushort32? is it accurate enough?



▶ Scaling and truncating the results from ushort → uchar

```
uchar32 V22, vout;
ushort32 vacc;
vout = (uchar32)((vacc+ (ushort32)V22 * 2 + (ushort32)V22 * 0) >> 5);
```

VPU0.vmac v3.uc32, #0x2, v34.s16, v35.s16, r7.uc, v14.c32



- Improvement suggestion
 - Each iteration includes:
 - 9 vectorized load instruction to load pixel data
 - 5 vmac3 instructions

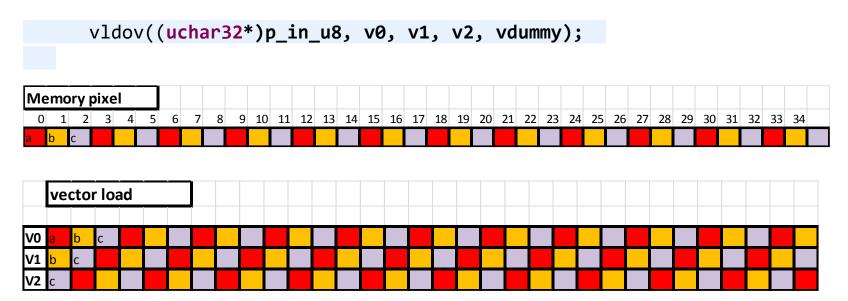


- optimization suggestion
 - Reduce the number of load slots to get better balance with the vmac3 operations
 - Use of vldov vec-c intrinsic

Gaussian3X3 load Operation Optimizations



- VIdov instruction
 - Can load up to 4 vectors with a predefined offset of memory location using both LSU0 and LSU1
 - We will replace 3 vector loads with one vldov operation.
 - ▶ Vldov in this case will load 3 vectors with 3 horizontal positions with offset of one pixel



Gaussian3X3 load Operation Optimizations



- Substitution of regular loads by VLDOV
- Ensure that CEVA_VLDOV definition is defined in gaussian3x3_lib.h and other definitions are not defined.

```
v00 = *(uchar32*)p in u8;
p in u8++;
v01 = *(uchar32*)p in u8;
p in u8++;
v02 = *(uchar32*)p in u8;
p in u8 += s32SrcStrideM2;
v10 = *(uchar32*)p in u8;
p in u8++;
v11 = *(uchar32*)p in u8;
p in u8++;
v12 = *(uchar32*)p in u8;
p in u8 += s32SrcStrideM2;
v20 = *(uchar32*)p in u8;
p in u8++;
v21 = *(uchar32*)p_in_u8;
p in u8++;
v22 = *(uchar32*)p in u8;
p in u8 += s32SrcStrideP2;
```



```
vldov((uchar32*)p_in_u8, v00, v01, v02, vdummy);
p_in_u8 += s32SrcStep;
vldov((uchar32*)p_in_u8, v10, v11, v12, vdummy);
p_in_u8 += s32SrcStep;
vldov((uchar32*)p_in_u8, v20, v21, v22, vdummy);
p_in_u8 -= s32SrcStep;
```

Why do we define the last vector as dummy variable vector?

Agenda



- Glossary
- Median Filter 3X3
 - Algorithm Analysis
 - Efficient VEC-C Code
 - Performance

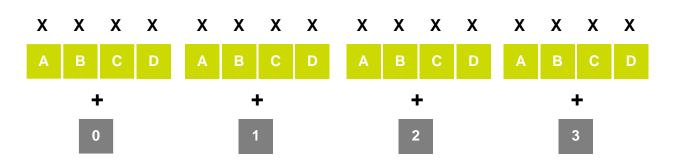
⋄Gaussian3X3

- Gaussian3X3 Kernel Overview
- Gaussian3X3 Vector Operators Code
- Gaussian3X3 load Operation Optimizations
- Gaussian3X3 Sliding window code



The vector instruction performs the arithmetic operation multiple times in parallel (producing eight, 16 or 32 results). The next sliding-window operation is offset by a step from the previous operation.







Loading vector coefficients

```
ushort coeff[16] = {2,4,2,0,4,8,4,0,2,4,2,0,0,0,0,0};
ushort16 v_coeff;
v_coeff = *(ushort16*)coeff;
```

Loading pixel data

```
uchar32 v00,V10,V20;

uchar* p_in_u8 = (uchar *)&p_u8Src[i*16];

v00 = *(uchar32*)p_in_u8;

p_in_u8+=s32SrcStep;

v10 = *(uchar32*)p_in_u8;

p_in_u8+=s32SrcStep;

v20 = *(uchar32*)p_in_u8;

p_in_u8+=s32SrcStep;
```



- Sliding window intrinsic
- We will use three different type of sliding window operations
 - Sliding window multiply
 - Sliding window multiple accumulate
 - Sliding window multiple accumulating with scaling



- Sliding window intrinsic
- Vswmpy5 serves also as initialization step of the accumulators
- uint16 vswmpy5(uchar32 inA, uchar32 inB, ushort16 inC, unsigned int inD);
 - ▶ InA 32 byte vector of loaded pixel data
 - InB 32 byte vector of loaded pixel data
 - InC 16 coefficients → in Gaussian3X3 only 9 will be non-zero.
 - inD offset and shift val control word
 - {inA, inB} X inC

Example

```
uchar32 v00;

ushort coeff[16] = {2,4,2,0, 4,8,4,0, 2,4,2,0, 0,0,0,0};

uint config0 = SW_CONFIG(0, 0, 0, 0, 0, 0);

uint16 vacc0;

vacc0 = (uint16)vswmpy5(v00, v00, v_coeff, config0);
```



- Sliding window intrinsic
- Vswmac5 multiple accumulating
- uint16 vswmac5(accumulate_sw, uchar32 inA, uchar32 inB, ushort16 inC, unsigned int inD, uint16 inoutW);
 - ▶ inA, InB, inC inA,InB pixel input vectors inC coefficient vector
 - ▶ inD input, coefficient offset and accumulation post shift
 - inoutW vector accumulator
 - accumulate_sw accumulation switch 'accumulate' is a predefined switch.

Example

```
uchar32 v10;

ushort coeff[16] = {2,4,2,0, 4,8,4,0, 2,4,2,0, 0,0,0,0};

uint config1 = SW_CONFIG(0, 0, 0, 4, 0, 0); // 4 for coeff offset

uint16 vacc0;

vacc0 = vswmac5(accumulate,v10, v10, v_coeff, config1, vacc0);
```

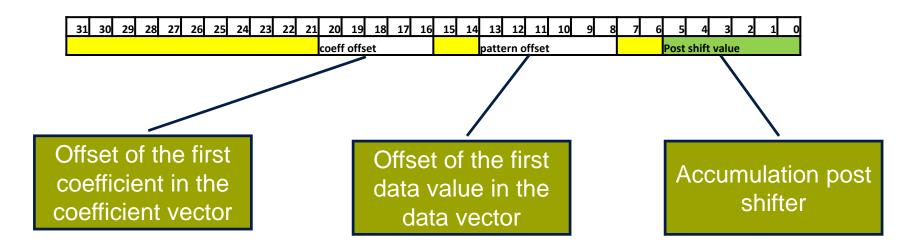


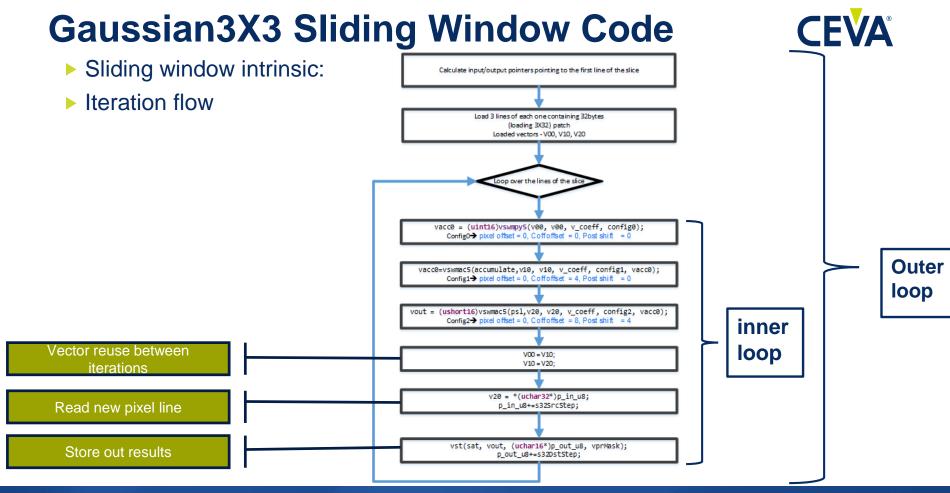
- Sliding window intrinsic:
- Vswmac5 accumulating and scaling
- short16 vswmac5(psl_sw, uchar32 inA, uchar32 inB, uchar32 inC, unsigned int inD, uint16 inE);
 - ▶ inA, InB, inC inA,InB pixel input vectors inC coefficient vector
 - ▶ inD − input, coefficient offset and accumulation post shift
 - ▶ inE vector accumulator
 - ▶ Psl_sw post shifter switch, the value 'psl' is a predefined value.
- Example

```
uchar32 v20;
Ushort16 vout;
ushort coeff[16] = {2,4,2,0, 4,8,4,0, 2,4,2,0, 0,0,0,0};
uint config2 = SW_CONFIG(5, 0, 0, 8, 0, 0); //5 for Post-shift (1/32)
vout= (ushort16)vswmac5(psl,v20, v20, v_coeff, config2, vacc0);
```



- Sliding window intrinsic:
- Control code field definition (configuration)
 - ► Control code is 32 bit width (assigned to one of the 32 scalar registers)







36

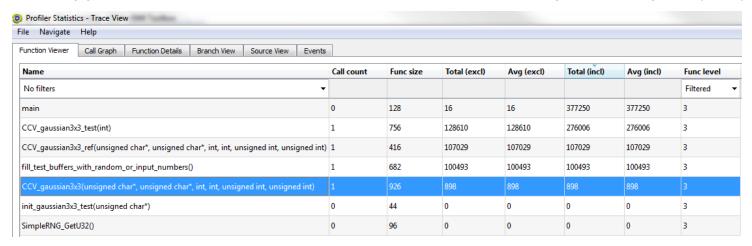
- Sliding window intrinsic:
- The algorithm will load 32 bytes on each iteration, but will store only 16
- Preparation between inner/outer loops:
 - Update pointer of current input/output line to point to the begging of the slice
 - ▶ Read 32 pixels for consecutive lines into 3 vectors



- Use the different slide window instruction as specified in slides 45-48
- Reuse loaded vectors each iteration in the inner loop should have only one load and one store.



▶ Run the application check bit-exact results and check the profiler report (ISS)



- Profile the same code using CAS Mode. What is the difference between ISS and CAS results?
- ▶ What is the theoretical cycles count? What is the compiler factor?



THANK YOU

www.ceva-dsp.com